Towards Practical First-Order Model Counting

Abstract

1 Introduction

Changes to the compilation process.

- We make the FCGs satisfy the following property. For every domain Δ, for every directed path P without Ref nodes, the number of atom counting or (generalised) domain recursion or independent partial grounding operations on Δ on P is at most one.

- when checking whether recursion is possible between a source and a target formula, the target formula (maybe the source formula as well?) must not include a domain as well as its subdomain. This might be equivalent to the formula being 'unshattered'. Such formulas could be seen as transitory, and allowing such recursive calls leads to FCGs that I cannot make sense of.

- During smoothing, when unit clauses (a.k.a. variables) are propagated in the opposite direction of the FCG arcs (i.e., 'upwards'), when visiting a `Ref` node, these clauses are translated using the domain map of the `Ref` node. For example, we might have a function g that depends on m (i.e., domain $a^\top$), which is called as $g(n)$ (i.e., on domain a). When smoothing that function call, we want to translate all mentions of $a^\top$ to a.

Proof of Correctness

1. Define what it means for a formula to be *consistent*.
   - The formula (including the mapping of constants to their domains) does not mention a domain together with its subdomain (i.e., all mentioned domains are pairwise disjoint). Specifically, we can't have a subdomain together with its parent domain or a subdomain together with a constant that belongs to its parent domain.
   - The target of `Ref` must also be a consistent formula.
   - The initial formula is consistent by definition.

2. Aim to show that Crane works correctly on all consistent formulas.

3. Prove that, whenever some compilation rule makes the formula ϕ inconsistent, greedy compilation rules will make ϕ consistent again before ϕ encounters non-greedy compilation rules that depend on ϕ being consistent.
   - In particular, prove that shattering and unit propagation will always eliminate the parent domain of the domains introduced by atom counting.
   - Constraint removal only applies when it can eliminate the parent domain entirely.
   - Prove this for a particular assumption about which compilation rules are set to be greedy. For the greedy algorithm, although all rules are applied in a greedy manner, we still assume that greedy rules will be applied before non-greedy rules.

The evaluation of base cases is done by simplifying the clauses and then using CRANE to find the base cases. First, while traversing the graph to find the equations, we store two maps: $\mathcal{F}$ (which stores the mapping from the function names to the formulae, whose model count they represent) and $\mathcal{D}$ (which stores the mapping from the variable names to the domains whose sizes they represent). Then, a particular domain is selected (using the algorithm described in previous reports), and the clauses are simplified. Then, CRANE is called on those clauses to evaluate the base cases. After that, we change the function names and variable to make it consistent with the previous domain to variable mapping, and append these base cases to the set of equations.

Contributions.

- Section 4
- Section 5
- Converting the recursive equations into a C++ program, which can then be compiled and executed to obtain numerical values (see Section 6).
- Support for infinite precision integers using the GNU Multiple Precision Arithmetic Library.
- The experiments are (going to be) more comprehensive than any other experimental study performed on (W)FOMC algorithms (TODO: references to my previous paper and some of the FORCLIFT and FASTWFOMC papers that contain experimental results).
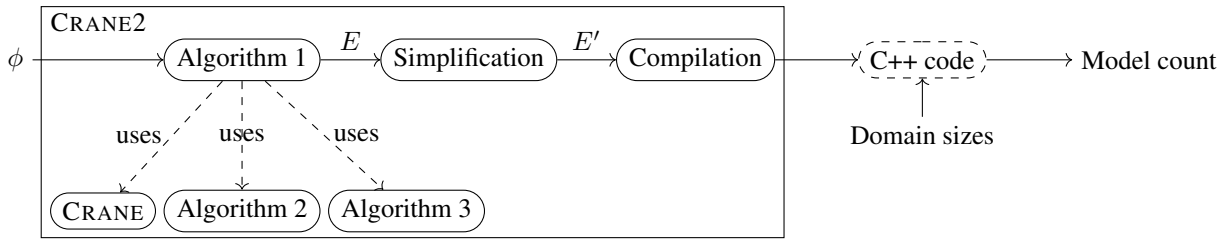
Figure 1: Using CRANE2 to compute the model count of formula ϕ. The formula is compiled into a set of equations E, which are then algebraically simplified and compiled to a C++ program. This program can then be run with different command line arguments to compute the model count of ϕ for various domain sizes.

Sections 4 and 6 deal with algebraic constructs whereas Section 5 deals with logic.

2   Preliminaries

2.1   Algebra

Notation. We write expr for an arbitrary algebraic expression. For any signature or clause C, argument or variable x, and number or constant t, we shall write $C[t/x]$ for the result of substituting t for all occurrences of x in C.

Definition 1. *A* function call *is a term of the form* $f(x_1 - c_1, \ldots, x_n - c_n)$ *(written* $f(\mathbf{x} - \mathbf{c})$ *for short), where* f *is an n-ary function, each x_i is a variable, and each c_i is a non-negative constant.*

Definition 2. *A* signature *is a term of the form* $f(x_1, \ldots, x_n)$ *(written* $f(\mathbf{x})$ *for short), where* f *is an n-ary function, and each x_i is a variable. The signature of a function call* $f(\mathbf{x} - \mathbf{c})$ *is* $f(\mathbf{x})$*. For example, the signature of* $f(x - 1, y - 2)$ *is* $f(x, y)$*.*

Definition 3. *An* equation *is always of the form* $f(\mathbf{x}) = \text{expr}$*, where* $f(\mathbf{x})$ *is a signature, and* expr *is an algebraic expression. Henceforth, we call* $f(\mathbf{x})$ *and* expr *the left-hand side (LHS) and the right-hand side (RHS) of the equation, respectively.*

Definition 4. *A* base case *is a function call where all arguments are either variables or constants (i.e., there is no subtraction within any argument), and there is at least one constant.*

2.2   Logic

A *formula* is a set of clauses. A *clause* is of the form $\forall x_1 \in \Delta_1 \forall x_2 \in \Delta_2 \ldots \forall x_n \in \Delta_n \phi(x_1, x_2, \ldots, x_n)$, where ϕ is a disjunction of literals that only contain variables $x_1, \ldots, x_n$ (and any constants). A *literal* is either an atom or its negation. An *atom* is of the form $P(t_1, \ldots, t_m)$, where P is a predicate, and $t_1, \ldots, t_m$ are terms. A *term* is either a variable or a constant.

For any clause C, let $\text{Preds}(C)$ denote the set of predicates in C. For any clause or literal l, let $\text{Vars}(l)$ denote the set of variables in l.

For any (bound) variable v, let $\text{Dom}(v)$ denote the domain over which v is quantified. For a literal l, $\text{Doms}(l) := \{ \text{Dom}(v) \mid v \in \text{Vars}(l) \}$. Similarly, for any clause C, let $\text{Doms}(C) = \{ \text{Dom}(v) \mid v \in \text{Vars}(C) \}$.

Algorithm 1: Finding the definitions of the base cases

Input: formula ϕ
Output: a set E of equations that define B
1 $(E, \mathcal{F}, \mathcal{D}) \leftarrow \text{Crane}(\phi)$;
2 **foreach** *base case* $f(\mathbf{x}) \in \text{FindBaseCases}(E)$ **do**
3 $\quad \psi \leftarrow \mathcal{F}(f)$;
4 $\quad$ **foreach** i *such that* x_i *is a constant* **do**
5 $\quad\quad \psi \leftarrow \text{Propagate}(\psi, \mathcal{D}(f, i), x_i)$;
6 $\quad (E', \_, \_) \leftarrow \text{Crane}(\psi)$;
7 $\quad E \leftarrow E \cup E'$;

Algorithm 2: $\text{FindBaseCases}(E)$

Input: set E of equations
Output: set B of base cases
1 $B \leftarrow \emptyset$;
2 **foreach** *equation* $(f(\mathbf{x}) = \text{expr}) \in E$ **do**
3 $\quad$ **foreach** *function call* $f(\mathbf{x} - \mathbf{c}) \in \text{expr}$ **do**
4 $\quad\quad$ **foreach** $c_i \in \mathbf{c}$ **do**
5 $\quad\quad\quad$ **for** $n \leftarrow 0$ **to** $c_i - 1$ **do**
6 $\quad\quad\quad\quad B \leftarrow B \cup \{ f(\mathbf{x})[x_i/n] \}$;

---

3   The Main Algorithm: Finding the Definitions of the Base Cases

See Algorithm 1

4   Identifying a Sufficient Set of Base Cases

We know that if, say, on the RHS of all equations, the domain size appears as $m - c_1, m - c_2, \ldots, m - c_k$, then finding $f(0, x_1, x_2, \ldots), f(1, x_1, x_2, \ldots), \ldots f(m_0, x_1, x_2, \ldots)$ for every function f, where $m_0 = \max(c_1, c_2, \ldots c_k) - 1$ forms a sufficient set of base cases. Hence, in order to do the same efficiently, we can take that domain for which m_0 is the minimum, i.e. $\text{argmin}(\max(c_1, c_2, \ldots c_k))$. Ideally, we should calculate the base cases by finding the base cases up to $\max(c_1, c_2, \ldots) - 1$. However, currently only empty and singleton domains are supported.

- line 3 iterates over all function calls in the algebraic expression.

- note that on **??** it is the signature and not the original function call.

The following steps were followed while finding the base cases:

First, expand the summations in each equation. Here we expand the summations of the form: $\sum_{x=0}^{x_1} \mathtt{expr} \cdot [a \le x < b]$ or similar inequalities where x is bounded by constants and a and b are constants, by substituting the value of x from a to $b-1$. For example, we replace $\sum_{x=0}^{x_1} \binom{x_1}{x} f(x_1 - x) \cdot [0 \le x < 2]$ by $\binom{x_1}{0} f(x_1) + \binom{x_1}{1} f(x_1 - 1)$.

Now, we find a domain that has only terms of the form $x - 1$ appearing on the RHS of the dependencies. The base cases are then calculated by setting this domain size to zero. For the above example, n is the selected domain, and not m since there are $m - 2$ and $m - 3$ terms appearing in the arguments.

The algorithm is described as Algorithm 2.

Example 1. *TODO: add an example where Algorithm 2 runs on either*

$$f_0(m,n) = f_1(m-1,n) + f_2(m,n-1)$$
$$f_1(m,n) = f_1(m-1,n-1) \times f_2(m-2,n-1)$$
$$f_2(m,n) = 2 \times f_1(m-3,n-1)$$

or

$$f(m,n) = g(m-1,n) + f(m-2,n-1)$$
$$g(m,n) = f(m-1,n-2) + g(m-1,n-1).$$

- *line 5: the limit is 2 for the arg $(x-3)$.*
- *line 6 for $f(\mathbf{x}) = f(y,z)$, $x_i = y$, and $n = 0$, add $f(y,z)[y/0] = f(0,z)$ to B.*

TODO: I changed the algorithm

Theorem 1. *Under the following assumptions, Algorithm 2 is guaranteed to return a sufficient set of base cases:*

- *there is no mutual recursion (i.e., cyclic dependencies between functions);*
- *each function f has exactly one equation with f on the LHS;*
- *in the recursive definition of function $f(x_1, \ldots, x_n)$, the i-th argument of each call to f on the RHS is of the form $x_i - c_i$, where:*
  - *$c_i \ge 0$ for all i, and*
  - *$c_i > 0$ for at least one i.*

Example 2. *For an equation $f(m,n) = 2 \times f(m-1,n)$, Algorithm 2 returns $f(0,n)$.*

NOTE: function calls such as $f(n-m)$ and $f(n-m-1)$ are ignored.

5 Propagating Domain Size Assumptions

5.1 Motivation: Why Tautologies Are Needed and Simply Removing The Clauses Does Not Work

Fact 1. *Assuming that domain Δ is empty, any clause that contains '$\forall x \in \Delta$' (for any variable x) is vacuously satisfied by all interpretations.*

Algorithm 3: `Propagate`(ϕ, Δ, n)

Input: formula ϕ, domain Δ, domain size $n \in \{0,1\}$
Output: formula ϕ'

1 $\phi' \leftarrow \emptyset$;
2 **if** $n = 0$ **then**
3   **foreach** *clause $C \in \phi$* **do**
4   **if** $\Delta \notin \mathrm{Doms}(C)$ **then** $\phi' \leftarrow \phi' \cup \{C\}$;
5        $C' \leftarrow \{l \in C \mid \Delta \notin \mathrm{Doms}(l)\}$;
6   **if** $\Delta \in \mathrm{Doms}(C)$ **and** $C' \ne \emptyset$ **then**
7           $l \leftarrow$ an arbitrary literal in C';
8           $\phi' \leftarrow \phi' \cup \{C' \cup \{\neg l\}\}$;
9 **else**
10     $c \leftarrow$ a new constant symbol;
11   **foreach** *clause $C \in \phi$* **do**
12        $C' \leftarrow C$;
13   **foreach** $v \in \mathrm{Vars}(C)$ *with* $\mathrm{Dom}(v) = \Delta$ **do**
14           $C' \leftarrow C'[c/v]$;
15        $\phi' \leftarrow \phi' \cup \{C'\}$;

For example, consider the formula

$$\forall x \in \Delta \forall y, z \in \Gamma : P(x) \vee Q(y,z) \qquad (1)$$
$$\forall y, z \in \Gamma' : Q(y,z) \qquad (2)$$

and assume that $\Gamma' \subseteq \Gamma$. In this case, if we set $|\Delta|$ to zero and remove clauses with variables quantified over Δ, we get

$$\forall y, z \in \Gamma' : Q(y,z), \qquad (3)$$

but the model count of Clause (3) is one. However, the actual model count should be $2^{|\Gamma|^2 - |\Gamma'|^2}$. That is, Q as a relation is a subset of $\Gamma \times \Gamma$. While Clause (1) becomes vacuously true, Clause (2) fixes the value of Q over $\Gamma' \times \Gamma' \subseteq \Gamma \times \Gamma$. Hence, the number of different values that Q can take is $|(\Gamma \times \Gamma) \setminus (\Gamma' \times \Gamma')| = |\Gamma|^2 - |\Gamma'|^2$.

We address this issue by converting clauses with universal quantifiers over the empty domain to tautologies, hence retaining all the predicates that have no argument assigned to the empty domain. For example, we would convert Clauses (1) and (2) to

$$\forall y, z \in \Gamma : Q(y,z) \vee \neg Q(y,z)$$
$$\forall y, z \in \Gamma' : Q(y,z).$$

The model count returned by this will also consider the truth value of Q over $y \in \Gamma \setminus \Gamma'$ or $z \in \Gamma \setminus \Gamma'$.

5.2 The Solution

We use Algorithm 3 to find the transformed formula corresponding to each base case obtained using Algorithm 2 and call CRANE on the formula to obtain the required base cases.

- Similar to generalised domain recursion (what's the difference?).
- Clauses where all literals have variables quantified over the empty domain are still removed.
- C' is guaranteed to be non-empty.

6   Generating C++ Code

The target is to generate C++ code that can evaluate numerical values of the model counts based on the equations generated by CRANE. We achieve this by parsing the equations generated by CRANE, simplifying them, and then generating C++ code. This approach can be done in linear time in the length of the formula using the Shunting Yard Algorithm.

The translation of a set E of equations into a C++ program works as follows.

First, we create a cache for each function in E. This is implemented as a multi-dimensional vector containing objects of `class cache_elem` defined as shown in the example code. The default initialization of this object is to -1 which is useful for recognizing unevaluated cases.

Next, we create a function definition for the LHS of each equation in E, including all functions and base cases. The signatures of these functions is decided as follows. A function call containing only variable arguments is named as the function itself, and ones with constants in their arguments are suffixed with a string that contains 'x' at the ith place if the ith argument is variable and the ith argument if that argument is a constant. For example, $f(x_1, x_2, x_3)$ is declared as `int f(int x1, int x2, int x3);` and $f(1, x_2, x_3)$ is declared as `int f_1xx(int x2, int x3);` (the constant arguments are removed from the signature).

The RHS of each equation in E is used to define the body of the equation corresponding to the LHS of that equation. The function body (for a function `func` corresponding to equation e) is formed as follows.

First, we check if the evaluation is already present in the cache. If so, then we return the cache element. The cache accesses are done using the `get_elem` function (definition given in the example), which resizes the cache if the accessed index is out of range.

Second, if the element is absent, then we decide if the arguments corresponding to e or one of the functions corresponding to the base cases, based on the value of the arguments. If it corresponds to the base cases, then we directly call the base case function and return its value. Else, we evaluate the value using the RHS, store the evaluated value in the cache and return the evaluated value. Note that in this step, we only call the base case function with one more constant argument than `func`. For example, `f0(x, y)` would call `f0_0x(y)` if $x = 0$ and `f0_x0(x)` if $y = 0$.

Third, to translate the RHS, we convert $\sum_{x=a}^{b}$ expr to

```
([y,z,...](){
    int sum = 0;
    for(int x = a; x <= b; x++)
        sum += expr;
    return sum;
})()
```

where $y, z, \ldots$ are the free variables present in expr.

7   Experiments

8   Conclusion

References