

Towards Practical First-Order Model Counting

Ananth K. Kidambi ✉

Indian Institute of Technology Bombay, Mumbai, India

Guramrit Singh ✉

Indian Institute of Technology Bombay, Mumbai, India

Paulius Dilkas ✉ 🏠 

University of Toronto, Toronto, Canada

Vector Institute, Toronto, Canada

Kuldeep S. Meel ✉ 🏠 

University of Toronto, Toronto, Canada

Abstract

First-order model counting (FOMC) is the problem of counting the number of models of a sentence in first-order logic. Since lifted inference techniques rely on reductions to variants of FOMC, the design of scalable methods for FOMC has attracted attention from both theoreticians and practitioners over the past decade. Recently, a new approach based on first-order knowledge compilation was proposed. This approach, called CRANE, instead of simply providing the final count, generates definitions of (possibly recursive) functions that can be evaluated with different arguments to compute the model count for any domain size. However, this approach is not fully automated, as it requires manual evaluation of the constructed functions. The primary contribution of this work is a fully automated compilation algorithm, called GANTRY, which transforms the function definitions into C++ code equipped with arbitrary-precision arithmetic. These additions allow the new FOMC algorithm to scale to domain sizes over 500,000 times larger than the current state of the art, as demonstrated through experimental results.

2012 ACM Subject Classification Theory of computation → Automated reasoning; Theory of computation → Logic and verification; Mathematics of computing → Combinatorics

Keywords and phrases first-order model counting, knowledge compilation, lifted inference

Funding *Ananth K. Kidambi*: TODO: funding

Guramrit Singh: TODO: funding

Paulius Dilkas: TODO: funding

Kuldeep S. Meel: TODO: funding

Acknowledgements Part of the research was conducted while all authors were at the National University of Singapore.

1 Introduction

First-order model counting (FOMC) is the task of determining the number of models for a sentence in first-order logic over a specified domain. The weighted variant, WFOMC, computes the total weight of these models, linking logical reasoning with probabilistic frameworks [32]. It builds upon earlier efforts in weighted model counting for propositional logic [3] and broader attempts to bridge logic and probability [15, 17, 20]. WFOMC is central to *lifted inference*, which enhances the efficiency of probabilistic calculations by exploiting symmetries [12]. Lifted inference continues to advance, with applications extending to constraint satisfaction problems [24] and probabilistic answer set programming [1]. Moreover, WFOMC has proven effective at reasoning over probabilistic databases [7] and probabilistic logic programs [18]. FOMC algorithms have also facilitated breakthroughs in discovering integer sequences [22] and developing recurrence relations for these sequences [5]. Recently,

these algorithms have been extended to perform sampling tasks [33].

The complexity of FOMC is generally measured by *data complexity*, with a formula classified as *liftable* if it can be solved in polynomial time relative to the domain size [10]. While all formulas with up to two variables are known to be liftable [29, 31], Beame et al. [2] demonstrated that liftability does not extend to all formulas, identifying an unliftable formula with three variables. Recent work has further extended the liftable fragment with additional axioms [23, 28] and counting quantifiers [13], expanding our understanding of liftability.

FOMC algorithms are diverse, with approaches ranging from *first-order knowledge compilation* (FOKC) to local search [16], Monte Carlo sampling [6], and anytime approximation [26]. Among these, FOKC-based algorithms are particularly prominent, transforming formulas into structured representations such as circuits or graphs. Notable examples include FORCLIFT [32] and its successor CRANE [5]. Another important algorithm, FASTWFOMC [27], uses cell enumeration as its foundation.

The CRANE algorithm marked a significant step forward, expanding the range of formulas handled by FOMC algorithms. However, it had notable limitations: it required manual evaluation of function definitions to compute model counts and introduced recursive functions without proper base cases, making it more complex to use. To address these shortcomings, we present GANTRY, a fully automated FOMC algorithm that overcomes the constraints of its predecessor. GANTRY can handle domain sizes over 500,000 times larger than previous algorithms and simplifies the user experience by automatically handling base cases and compiling function definitions into efficient C++ programs.

In Section 2, we cover some preliminaries, and in Section 3, we detail all our technical contributions. Finally, in Section 4, we present our experimental results, demonstrating GANTRY’s performance compared to other FOMC algorithms, and, in Section 5, we conclude the paper by discussing promising avenues for future work.

2 Preliminaries

In Section 2.1, we summarise the basic principles of first-order logic. Then, in Section 2.2, we formally define (W)FOMC and discuss the distinctions between three variations of first-order logic used for FOMC. Finally, in Section 2.3, we introduce the terminology used to describe the output of the original CRANE algorithm, i.e., functions and equations that define them.

We use \mathbb{N}_0 to represent the set of non-negative integers. In both algebra and logic, we write $S\sigma$ to denote the application of a *substitution* σ to an expression S , where $\sigma = [x_1 \mapsto y_1, x_2 \mapsto y_2, \dots, x_n \mapsto y_n]$ signifies the replacement of all instances of x_i with y_i for all $i = 1, \dots, n$.

2.1 First-Order Logic

In this section, we will review the basic concepts of first-order logic as they are used in FOKC algorithms. There are two key differences between the logic used by these algorithms and the logic supported as input. First, Skolemization [31] eliminates existential quantifiers by introducing additional predicates. Please note that Skolemization here differs from the standard Skolemization procedure that introduces function symbols [9]. Second, the input formula is rewritten as a conjunction of clauses, each in *prenex normal form* [8].

A *term* can be either a variable or a constant. An *atom* can be either $P(t_1, \dots, t_m)$ (or $P(\mathbf{t})$) for some predicate P and terms t_1, \dots, t_m or $x = y$ for some terms x and y . The *arity* of a predicate is the number of arguments it takes, i.e., m in the case of the predicate P mentioned above. We write P/m to denote a predicate along with its arity. A *literal*

Logic	Sorts	Constants	Variables	Quantifiers	Additional atoms
FO	one or more	✓	unlimited	\forall, \exists	$x = y$
C^2	one	✗	two	$\forall, \exists, \exists^{=k}, \exists^{\leq k}, \exists^{\geq k}$	—
$UFO^2 + CC$	one	✗	two	\forall	$ P = m$

■ **Table 1** A comparison of the three logics used in FOMC. The 2nd–5th columns refer to: the number of sorts, support for constants, the maximum number of variables, and supported quantifiers, respectively. The last column lists supported atoms in addition to those of the form $P(\mathbf{t})$ for a predicate P/n and an n -tuple of terms \mathbf{t} . Here: k and m are non-negative integers, with the latter depending on the domain size, P represents a predicate, and x and y are terms.

can be either an atom (i.e., a *positive* literal) or its negation (i.e., a *negative* literal). An atom is *ground* if it contains no variables, i.e., only constants. A *clause* is of the form $\forall x_1 \in \Delta_1. \forall x_2 \in \Delta_2 \dots \forall x_n \in \Delta_n. \phi(x_1, x_2, \dots, x_n)$, where ϕ is a disjunction of literals that only contain variables x_1, \dots, x_n (and any constants). We say that a clause is a (*positive*) *unit clause* if there is only one literal with a predicate, and it is a positive literal. Finally, a *formula* is a conjunction of clauses. Throughout the paper, we will use set-theoretic notation, interpreting a formula as a set of clauses and a clause as a set of literals.

2.2 FOMC Algorithms and Their Logics

In Table 1, we outline the differences among three first-order logics commonly used in FOMC: FO, C^2 , and $UFO^2 + CC$. First, FO is the input format for FORCLIFT* and its extensions CRANE† and GANTRY. Second, C^2 is often used in the literature on FASTWFOMC and related methods [13, 14]. Finally, $UFO^2 + CC$ is the input format supported by the most recent implementation of FASTWFOMC‡. The notation we use to refer to each logic is standard in the case of C^2 and $UFO^2 + CC$ [25] and redefined to be more specific in the case of FO. All three logics are function-free, and domains are always assumed to be finite. As usual, we presuppose the *unique name assumption*, which states that two constants are equal if and only if they are the same constant [19].

In FO, each term is assigned to a *sort*, and each predicate P/n is assigned to a sequence of n sorts. Each sort has its corresponding domain. These assignments to sorts are typically left implicit and can be reconstructed from the quantifiers. For example, $\forall x, y \in \Delta. P(x, y)$ implies that variables x and y have the same sort. On the other hand, $\forall x \in \Delta. \forall y \in \Gamma. P(x, y)$ implies that x and y have different sorts, and it would be improper to write, for example, $\forall x \in \Delta. \forall y \in \Gamma. P(x, y) \vee x = y$. FO is also the only logic to support constants, formulas with more than two variables, and the equality predicate. While we do not explicitly refer to sorts in subsequent sections of this paper, the many-sorted nature of FO is paramount to the algorithms presented therein.

► **Remark 1.** In the case of FORCLIFT and its extensions, support for a formula as valid input does not imply that the algorithm can compile the formula into a circuit or graph suitable for lifted model counting. However, it is known that FORCLIFT compilation is guaranteed to succeed on any FO formula without constants and with at most two variables [29, 31].

* <https://github.com/UCLA-StarAI/Forclift>

† <https://doi.org/10.5281/zenodo.8004077>

‡ <https://github.com/jan-toth/FastWFOMC.jl>

Compared to FO, C^2 and $UFO^2 + CC$ lack support for constants, the equality predicate, multiple domains, and formulas with more than two variables. The advantage that C^2 brings over FO is the inclusion of *counting quantifiers*. That is, alongside \forall and \exists , C^2 supports $\exists^{=k}$, $\exists^{\leq k}$, and $\exists^{\geq k}$ for any positive integer k . For example, $\exists^{=1}x. \phi(x)$ means that there exists *exactly one* x such that $\phi(x)$, and $\exists^{\leq 2}x. \phi(x)$ means that there exist *at most two* such x . $UFO^2 + CC$, on the other hand, does not support any existential quantifiers but instead incorporates (*equality*) *cardinality constraints*. For example, $|P| = 3$ constrains all models to have *precisely three positive literals with the predicate P* .

► **Definition 2 (Model).** Let ϕ be a formula in FO. For each predicate P/n in ϕ , let $(\Delta_i^P)_{i=1}^n$ be a list of the corresponding domains. Let σ be a map from the domains of ϕ to their interpretations as sets such that the sets are pairwise disjoint, and the constants in ϕ are included in the corresponding domains. A structure of ϕ is a set M of ground literals defined by adding to M either $P(\mathbf{t})$ or $\neg P(\mathbf{t})$ for every predicate P/n in ϕ and n -tuple $\mathbf{t} \in \prod_{i=1}^n \sigma(\Delta_i^P)$. A structure is a model if it satisfies ϕ .

► **Remark 3.** The distinctness of domains is important in two ways. First, in terms of expressiveness, a clause such as $\forall x \in \Delta. P(x, x)$ is valid if predicate P is defined over two copies of the same domain and invalid otherwise. Second, having more distinct domains makes the problem more decomposable for the FOKC algorithm. With distinct domains, the algorithm can make assumptions or deductions about, e.g., the first domain of predicate P without worrying how (or if) they apply to the second domain.

While this work focuses on FOMC, we still define the weighted variant of the problem as Skolemization relies on weights even for unweighted FOMC.

► **Definition 4 (WFOMC instance).** A WFOMC instance comprises: a formula ϕ in FO, two (rational) weights $w^+(P)$ and $w^-(P)$ assigned to each predicate P in ϕ , and σ as described in Definition 2. Unless specified otherwise, we assume all weights to be equal to 1.

► **Definition 5 (WFOMC [32]).** Given a WFOMC instance (ϕ, w^+, w^-, σ) as in Definition 4, the (symmetric) weighted first-order model count (WFOMC) of ϕ is

$$\sum_{M \models \phi} \prod_{P(\mathbf{t}) \in M} w^+(P) \prod_{\neg P(\mathbf{t}) \in M} w^-(P), \quad (1)$$

where the sum is over all models of ϕ .

► **Example 6 (Counting functions).** To define predicate P as a function from a domain Δ to itself, in C^2 one would write $\forall x \in \Delta. \exists^{=1} y \in \Delta. P(x, y)$. In $UFO^2 + CC$, the same could be written as


$$(\forall x, y \in \Delta. S(x) \vee \neg P(x, y)) \wedge (|P| = |\Delta|), \quad (2)$$

where $w^-(S) = -1$. Although Formula (2) has more models compared to its counterpart in C^2 , the negative weight $w^-(S) = -1$ makes some of the terms in Equation (1) cancel out.

Equivalently, in FO we would write

$$(\forall x \in \Gamma. \exists y \in \Delta. P(x, y)) \wedge (\forall x \in \Gamma. \forall y, z \in \Delta. P(x, y) \wedge P(x, z) \Rightarrow y = z). \quad (3)$$

The first clause asserts that each x must have at least one corresponding y , while the second statement adds the condition that if x is mapped to both y and z , then y must equal z . It is important to note that Formula (3) is written with two domains instead of just one. However, we can still determine the correct number of functions by assuming that the sizes of Γ and Δ are equal. This formulation, as observed by Dilkas and Belle [5], can prove beneficial in enabling FOKC algorithms to find efficient solutions.



■ **Figure 1** The outline of using GANTRY to compute the model count of a formula ϕ . First, the formula is compiled into a set of equations, which are then used to create a C++ program. This program can be executed with different command line arguments to calculate the model count of ϕ for different domain sizes. To accomplish this, the `CompileWithBaseCases` function makes use of the FOKC algorithm of CRANE, algebraic simplification techniques (denoted as `Simplify`), and two other auxiliary procedures.

2.3 Algebra

We write `expr` to represent an arbitrary algebraic expression. It is important to note that some terms have different meanings in algebra and logic. In algebra, a *constant* refers to a non-negative integer. Likewise, a *variable* can either be a parameter of a function or a variable introduced through summation, such as i in the expression $\sum_{i=1}^n \text{expr}$. A (function) *signature* is $f(x_1, \dots, x_n)$ (or $f(\mathbf{x})$ for short), where f represents an n -ary function, and each x_i represents a variable. An *equation* is $f(\mathbf{x}) = \text{expr}$, with $f(\mathbf{x})$ representing a signature.

► **Definition 7** (Base case). Let $f(\mathbf{x})$ be a function call where each x_i is either a constant or a variable (note that signatures are included in this definition). Then function call $f(\mathbf{y})$ is considered a base case of $f(\mathbf{x})$ if $f(\mathbf{y}) = f(\mathbf{x})\sigma$, where σ is a substitution that replaces one or more x_i with a constant.

3 Technical Contributions

Figure 1 provides an overview of GANTRY’s workflow. Section 3.1 describes the main algorithm for completing the definitions of recursive functions with a sufficient set of base cases. Sections 3.2 and 3.3 describe subsidiary algorithms for constructing a set of base cases and their corresponding logical formulas. Section 3.4 explains the post-processing techniques for ensuring accurate model counting. Additionally, Section 3.5 explains the process of compiling function definitions into C++ code, greatly expanding upon the range of formulas that could previously be handled by similar approaches [11].

3.1 Completing the Definitions of Functions

Before describing the main contribution of this work, let us review the essential aspects of FOKC as realised by CRANE. The input formula is compiled into: set \mathcal{E} of equations, map \mathcal{F} from function names to formulas, and map \mathcal{D} from function names and argument indices to domains. \mathcal{E} can contain any number of functions, one of which (denoted by f) represents the solution to the FOMC problem. To compute the FOMC for particular domain sizes, f must be evaluated with those domain sizes as arguments. \mathcal{D} records this correspondence between function arguments and domains.

Algorithm 1 `CompileWithBaseCases(ϕ)

---`

Input: formula ϕ
Output: set \mathcal{E} of equations
1 $(\mathcal{E}, \mathcal{F}, \mathcal{D}) \leftarrow \text{CRANE}(\phi);$
2 $\mathcal{E} \leftarrow \text{Simplify}(\mathcal{E});$
3 **foreach** *base case* $f(\mathbf{x}) \in \text{FindBaseCases}(\mathcal{E})$ **do**
4 $\psi \leftarrow \mathcal{F}(f);$
5 **foreach** *index* i *such that* $x_i \in \mathbb{N}_0$ **do** $\psi \leftarrow \text{Propagate}(\psi, \mathcal{D}(f, i), x_i);$
6 $\mathcal{E} \leftarrow \mathcal{E} \cup \text{CompileWithBaseCases}(\psi);$

190 Algorithm 1 presents our overall approach for compiling a formula into equations that
191 include the necessary base cases. To begin, we use the FOKC algorithm of the original
192 CRANE to compile the formula into the three components: \mathcal{E} , \mathcal{F} , and \mathcal{D} . After some algebraic
193 simplification, \mathcal{E} is passed to the `FindBaseCases` procedure (see Section 3.2). For each base
194 case $f(\mathbf{x})$, we retrieve the logical formula $\mathcal{F}(f)$ associated with the function name f and
195 simplify it using the `Propagate` procedure (explained in detail in Section 3.3). We do this by
196 iterating over all indices of \mathbf{x} , where x_i is a constant, and using `Propagate` to simplify ψ by
197 assuming that domain $\mathcal{D}(f, i)$ has size x_i . Finally, on line 6, `CompileWithBaseCases` recurses
198 on these simplified formulas and adds the resulting base case equations to \mathcal{E} . Example 9
199 below provides more detail.

200 ► **Remark 8.** Although `CompileWithBaseCases` starts with a call to `CRANE`, the proposed
201 algorithm is not just a post-processing step for FOKC because Algorithm 1 is recursive and
202 can issue more calls to `CRANE` on various derived formulas.

203 ► **Example 9 (Counting bijections).** Consider the following formula (previously examined by
204 Dilkas and Belle [5]) that defines predicate P as a bijection between two sets Γ and Δ :

$$\begin{aligned}
& (\forall x \in \Gamma. \exists y \in \Delta. P(x, y)) \wedge \\
& (\forall y \in \Delta. \exists x \in \Gamma. P(x, y)) \wedge \\
205 \quad & (\forall x \in \Gamma. \forall y, z \in \Delta. P(x, y) \wedge P(x, z) \Rightarrow y = z) \wedge \\
& (\forall x, z \in \Gamma. \forall y \in \Delta. P(x, y) \wedge P(z, y) \Rightarrow x = z).
\end{aligned}$$

206 We specifically examine the first solution returned by GANTRY for this formula.

207 After line 2, we have

$$\begin{aligned}
208 \quad \mathcal{E} &= \left\{ \begin{array}{l} f(m, n) = \sum_{l=0}^n \binom{n}{l} (-1)^{n-l} g(l, m), \\ g(l, m) = g(l-1, m) + mg(l-1, m-1) \end{array} \right\}; \\
209 \quad \mathcal{D} &= \{ (f, 1) \mapsto \Gamma, (f, 2) \mapsto \Delta, (g, 1) \mapsto \Delta^\top, (g, 2) \mapsto \Gamma \},
\end{aligned}$$

210 where Δ^\top is a new domain. (We omit the definition of \mathcal{F} as the formulas can get a bit
211 verbose.) Then `FindBaseCases` identifies two base cases: $g(0, m)$ and $g(l, 0)$. In both cases,
212 `CompileWithBaseCases` recurses on the formula $\mathcal{F}(g)$ simplified by assuming that one of the
213 domains is empty. In the first case, we recurse on the formula $\forall x \in \Gamma. S(x) \vee \neg S(x)$, where
214 S is a predicate introduced by Skolemization with weights $w^+(S) = 1$ and $w^-(S) = -1$.
215 Hence, we obtain the base case $g(0, m) = 0^m$. In the case of $g(l, 0)$, `Propagate`($\psi, \Gamma, 0$)
216 returns an empty formula, resulting in $g(l, 0) = 1$.

Algorithm 2 FindBaseCases(\mathcal{E})

Input: set \mathcal{E} of equations**Output:** set \mathcal{B} of base cases

```

1  $\mathcal{B} \leftarrow \emptyset$ ;
2 foreach function call  $f(\mathbf{y})$  on the right-hand side of an equation in  $\mathcal{E}$  do
3    $\mathbf{x} \leftarrow$  the parameters of  $f$  in its definition;
4   foreach  $y_i \in \mathbf{y}$  do
5     if  $y_i \in \mathbb{N}_0$  then  $\mathcal{B} \leftarrow \mathcal{B} \cup \{f(\mathbf{x})[x_i \mapsto y_i]\}$ ;
6     else if  $y_i = x_i - c_i$  for some  $c_i \in \mathbb{N}_0$  then
7       for  $j \leftarrow 0$  to  $c_i - 1$  do  $\mathcal{B} \leftarrow \mathcal{B} \cup \{f(\mathbf{x})[x_i \mapsto j]\}$ ;

```

It is worth noting that these base cases overlap when $l = m = 0$ but remain consistent since $0^0 = 1$. Generally, let ϕ be a formula with two domains Γ and Δ , and let $n, m \in \mathbb{N}_0$. Then the FOMC of $\text{Propagate}(\phi, \Delta, n)$ assuming $|\Gamma| = m$ is the same as the FOMC of $\text{Propagate}(\phi, \Gamma, m)$ assuming $|\Delta| = n$.

Finally, the main responsibility of the **Simplify** procedure is to handle the algebraic pattern $\sum_{m=0}^n [a \leq m \leq b] f(m)$. Here: n is a variable, $a, b \in \mathbb{N}_0$ are constants, and f is an expression that may depend on m . Additionally, $[a \leq m \leq b] = \begin{cases} 1 & \text{if } a \leq m \leq b \\ 0 & \text{otherwise} \end{cases}$. **Simplify** transforms this pattern into $f(a) + f(a+1) + \dots + f(\min\{n, b\})$. For instance, in the case of Example 9, **Simplify** transforms $g(l, m) = \sum_{k=0}^m [0 \leq k \leq 1] \binom{m}{k} g(l-1, m-k)$ into $g(l, m) = g(l-1, m) + mg(l-1, m-1)$.

3.2 Identifying a Sufficient Set of Base Cases

Algorithm 2 summarises the implementation of **FindBaseCases**. It considers two types of arguments when a function f calls itself recursively: constants and arguments of the form $x_i - c_i$. Here, c_i is a constant, and x_i is the i -th argument of the signature of f . When the argument is a constant c_i , a base case with c_i is added. In the second case, a base case is added for each constant from 0 up to (but not including) c_i .

► **Example 10.** Consider the recursive function g from Example 9. **FindBaseCases** iterates over two function calls: $g(l-1, m)$ and $g(l-1, m-1)$. The former produces the base case $g(0, m)$, while the latter produces both $g(0, m)$ and $g(l, 0)$.

It can be shown that the base cases identified by **FindBaseCases** are sufficient for the algorithm to terminate.⁴

► **Theorem 11 (Termination).** *Let \mathcal{E} represent the equations returned by **CompileWithBaseCases**. Let f be an n -ary function in \mathcal{E} and $\mathbf{x} \in \mathbb{N}_0^n$. Then the evaluation of $f(\mathbf{x})$ terminates.*

We prove Theorem 11 using double induction. First, we apply induction to the number of functions in \mathcal{E} . Then, we use induction on the arity of the ‘last’ function in \mathcal{E} according to some topological ordering. For the detailed proof, please refer to the technical appendix.

⁴ Note that characterising the fine-grained complexity of the solutions found by GANTRY or other FOMC algorithms is an emerging area of research. These questions have been partially addressed in previous work [5, 25] and are orthogonal to the goals of this section.

■ **Algorithm 3** $\text{Propagate}(\phi, \Delta, n)$

Input: formula ϕ , domain Δ , $n \in \mathbb{N}_0$
Output: formula ϕ'

```

1   $\phi' \leftarrow \emptyset$ ;
2  if  $n = 0$  then
3      foreach clause  $C \in \phi$  do
4          if  $\Delta \notin \text{Doms}(C)$  then  $\phi' \leftarrow \phi' \cup \{C\}$ ;
5          else
6               $C' \leftarrow \{l \in C \mid \Delta \notin \text{Doms}(l)\}$ ;
7              if  $C' \neq \emptyset$  then
8                   $l \leftarrow$  an arbitrary literal in  $C'$ ;
9                   $\phi' \leftarrow \phi' \cup \{C' \cup \{\neg l\}\}$ ;
10 else
11      $D \leftarrow$  a set of  $n$  new constants in  $\Delta$ ;
12     foreach clause  $C \in \phi$  do
13          $(x_i)_{i=1}^m \leftarrow$  the variables in  $C$  with domain  $\Delta$ ;
14         if  $m = 0$  then  $\phi' \leftarrow \phi' \cup \{C\}$ ;
15         else  $\phi' \leftarrow \phi' \cup \{C[x_1 \mapsto c_1, \dots, x_m \mapsto c_m] \mid (c_i)_{i=1}^m \in D^m\}$ ;

```

243 3.3 Propagating Domain Size Assumptions

244 Algorithm 3, called **Propagate**, modifies the formula ϕ based on the assumption that $|\Delta| = n$.
245 When $n = 0$, some clauses become vacuously satisfied and can be removed. When $n > 0$,
246 partial grounding is performed by replacing all variables quantified over Δ with constants.
247 (None of the formulas examined in this work had $n > 1$.) Algorithm 3 handles these two
248 cases separately. For a literal or a clause C , the set of corresponding domains is denoted as
249 $\text{Doms}(C)$.

250 In the case of $n = 0$, there are three types of clauses to consider:

- 251 1. those that do not mention Δ ,
- 252 2. those in which every literal contains variables quantified over Δ , and
- 253 3. those that have some literals with variables quantified over Δ and some without.

254 Clauses of Type 1 are transferred to the new formula ϕ' without any changes. For clauses of
255 Type 2, C' is empty, so these clauses are filtered out. As for clauses of Type 3, a new kind of
256 smoothing is performed, which will be explained in Section 3.4.

257 In the case of $n > 0$, n new constants are introduced. Let C be an arbitrary clause in ϕ ,
258 and let $m \in \mathbb{N}_0$ be the number of variables in C quantified over Δ . If $m = 0$, C is added
259 directly to ϕ' . Otherwise, a clause is added to ϕ' for every possible combination of replacing
260 the m variables in C with the n new constants.

261 ► **Example 12.** Let $C \equiv \forall x \in \Gamma. \forall y, z \in \Delta. \neg P(x, y) \vee \neg P(x, z) \vee y = z$. Then $\text{Doms}(C) =$
262 $\text{Doms}(\neg P(x, y)) = \text{Doms}(\neg P(x, z)) = \{\Gamma, \Delta\}$, and $\text{Doms}(y = z) = \{\Delta\}$. A call to

263 `Propagate`($\{C\}, \Delta, 3$) would result in the following formula with nine clauses:

$$\begin{aligned}
 264 \quad & (\forall x \in \Gamma. \neg P(x, c_1) \vee \neg P(x, c_1) \vee c_1 = c_1) \wedge \\
 265 \quad & (\forall x \in \Gamma. \neg P(x, c_1) \vee \neg P(x, c_2) \vee c_1 = c_2) \wedge \\
 266 \quad & \vdots \\
 267 \quad & (\forall x \in \Gamma. \neg P(x, c_3) \vee \neg P(x, c_3) \vee c_3 = c_3).
 \end{aligned}$$

268 Here, c_1 , c_2 , and c_3 are the new constants.

269 3.4 Smoothing the Base Cases

270 *Smoothing* modifies a circuit to reintroduce eliminated atoms, ensuring the correct model
 271 count [4, 32]. In this section, we describe a similar process performed on lines 7–9 of
 272 Algorithm 3. Line 7 checks if smoothing is necessary, and lines 8 and 9 execute it. If the
 273 condition on line 7 is not satisfied, the clause is not smoothed but omitted.

274 Suppose `Propagate` is called with arguments $(\phi, \Delta, 0)$, i.e., we are simplifying the formula
 275 ϕ by assuming that the domain Δ is empty. Informally, if there is a predicate P in ϕ unrelated
 276 to Δ , smoothing preserves all occurrences of P even if all clauses with P become vacuously
 277 satisfied.

278 ► **Example 13.** Let ϕ be

$$279 \quad (\forall x \in \Delta. \forall y, z \in \Gamma. Q(x) \vee P(y, z)) \wedge \quad (4)$$

$$280 \quad (\forall y, z \in \Gamma'. P(y, z)), \quad (5)$$

281 where $\Gamma' \subseteq \Gamma$ is a domain introduced by a compilation rule. It should be noted that P , as a
 282 relation, is a subset of $\Gamma \times \Gamma$.

283 Now, let us reason manually about the model count of ϕ when $\Delta = \emptyset$. Predicate Q can
 284 only take one value, $Q = \emptyset$. The value of P is fixed over $\Gamma' \times \Gamma'$ by Clause (5), but it can vary
 285 freely over $(\Gamma \times \Gamma) \setminus (\Gamma' \times \Gamma')$ since Clause (4) is vacuously satisfied by all structures. Therefore,
 286 the correct FOMC should be $2^{|\Gamma|^2 - |\Gamma'|^2}$. However, without line 9, `Propagate` would simplify
 287 ϕ to $\forall y, z \in \Gamma'. P(y, z)$. In this case, P is a subset of $\Gamma' \times \Gamma'$. This simplified formula has
 288 only one model: $\{P(y, z) \mid y, z \in \Gamma'\}$. By including line 9, `Propagate` transforms ϕ to

$$289 \quad (\forall y, z \in \Gamma. P(y, z) \vee \neg P(y, z)) \wedge (\forall y, z \in \Gamma'. P(y, z)),$$

290 which retains the correct model count.

291 It is worth mentioning that the choice of l on line 8 of Algorithm 3 is inconsequential
 292 because any choice achieves the same goal: constructing a tautological clause that retains
 293 the literals in C' .

294 3.5 Generating C++ Code

295 In this section, we will describe the final step of GANTRY as outlined in Figure 1. This step
 296 involves translating the set of equations \mathcal{E} into C++ code. The resulting C++ program
 297 can then be compiled and executed with different command-line arguments to compute the
 298 model count of the formula for various domain sizes.

299 Each equation in \mathcal{E} is compiled into a C++ function, along with a separate cache for
 300 memoisation. Let us consider an arbitrary equation $e = (f(\mathbf{x}) = \text{expr}) \in \mathcal{E}$, and let $\mathbf{c} \in \mathbb{N}_0^n$

represent the arguments of the corresponding C++ function. The implementation of e consists of three parts. First, we check if \mathbf{c} is already present in the cache of e . If it is, we simply return the cached value. Second, for each base case $f(\mathbf{y})$ of $f(\mathbf{x})$ (as defined in Definition 7), we check if \mathbf{c} *matches* \mathbf{y} , i.e., $c_i = y_i$ whenever $y_i \in \mathbb{N}_0$. If this condition is satisfied, \mathbf{c} is redirected to the C++ function that corresponds to the definition of the base case $f(\mathbf{y})$. Finally, if none of the above cases apply, we evaluate \mathbf{c} based on the expression expr , store the result in the cache, and return it.

4 Experimental Evaluation

Our empirical evaluation sought to compare the runtime performance of GANTRY with the current state of the art, namely FASTWFOMC and FORCLIFT. It is worth remarking that FORCLIFT does not support arbitrary precision, and returns error for cases that requires arbitrary precision reasoning. Our experiments involve two versions of GANTRY: GANTRY-GREEDY and GANTRY-BFS. Like its predecessor, GANTRY has two modes for applying compilation rules to formulas: one that uses a greedy search algorithm similar to FORCLIFT and another that combines greedy and breadth-first search.

The experiments were conducted using an Intel Skylake 2.4 GHz CPU with 188 GiB of memory and CentOS 7. C++ programs were compiled using the Intel C++ Compiler 2020u4. FASTWFOMC ran on Julia 1.10.4, while the other algorithms were executed on the Java Virtual Machine 1.8.0_201.

4.1 Benchmarks

We compare these algorithms using three benchmarks from previous studies. The first benchmark is the function-counting problem from Example 6, previously examined by Dilkas and Belle [5]. The second benchmark is a variant of the well-known ‘Friends and Smokers’ Markov logic network [21, 30]. In \mathcal{C}^2 , FO, and $\text{UFO}^2 + \text{CC}$, this problem can be formulated as

$$(\forall x, y \in \Delta. S(x) \wedge F(x, y) \Rightarrow S(y)) \wedge (\forall x \in \Delta. S(x) \Rightarrow C(x))$$

or, equivalently, in conjunctive normal form as

$$(\forall x, y \in \Delta. S(y) \vee \neg S(x) \vee \neg F(x, y)) \wedge (\forall x \in \Delta. C(x) \vee \neg S(x)).$$

Finally, we include the bijection-counting problem previously utilised by Dilkas and Belle [5]. Its formulation in FO is described in Example 9. The equivalent formula in \mathcal{C}^2 is

$$(\forall x \in \Delta. \exists^{=1} y \in \Delta. P(x, y)) \wedge (\forall y \in \Delta. \exists^{=1} x \in \Delta. P(x, y)).$$

Similarly, in $\text{UFO}^2 + \text{CC}$ the same formula can be written as

$$(\forall x, y \in \Delta. R(x) \vee \neg P(x, y)) \wedge (\forall x, y \in \Delta. S(x) \vee \neg P(y, x)) \wedge (|P| = |\Delta|),$$

where $w^-(R) = w^-(S) = -1$.

The three benchmark families cover a wide range of possibilities. The ‘friends’ benchmark stands out as it uses multiple predicates and can be expressed in FO using just two variables without cardinality constraints or counting quantifiers. The ‘functions’ benchmark, on the other hand, can still be handled by all the algorithms, but it requires cardinality constraints, counting quantifiers, or more than two variables. Lastly, the ‘bijections’ benchmark is an example of a formula that FASTWFOMC can handle but FORCLIFT cannot.




Figure 2 The runtime of the algorithms as a function of the domain size. Note that both axes are on a logarithmic scale.

For evaluation purposes, we ran each algorithm on each benchmark using domains of sizes $2^1, 2^2, 2^3$, and so on, until an algorithm failed to handle a domain size due to timeout, out of memory error, or out of precision errors. While we separately measured compilation and inference time, we primarily focus on total runtime, dominated by the latter.

4.2 Results

Figure 2 presents a summary of the experimental results. Only FASTWFOMC and GANTRY-BFS could handle the bijection-counting problem. For this benchmark, the largest domain sizes these algorithms could accommodate were 64 and 4096, respectively. On the other two benchmarks, FORCLIFT had the lowest runtime. However, due to its finite precision, it only scaled up to domain sizes of 16 and 128 for ‘friends’ and ‘functions’, respectively. FASTWFOMC outperformed FORCLIFT in the case of ‘friends’, but not ‘functions’, as it could handle domains of size 1024 and 64, respectively. Furthermore, both GANTRY-BFS and GANTRY-GREEDY performed similarly on both benchmarks. Similarly to the ‘bijections’ benchmark, GANTRY significantly outperformed the other two algorithms, scaling up to domains of size 8192 and 67,108,864, respectively.

Another aspect of the experimental results that deserves separate discussion is compilation. Both Julia and Scala use just-in-time (JIT) compilation, which means that FASTWFOMC and FORCLIFT take longer to run on the smallest domain size, where most JIT compilation occurs. In the case of GANTRY, it is only run once per benchmark, so the JIT compilation time is included in its overall runtime across all domain sizes. Additionally, while FORCLIFT’s compilation is generally faster than that of GANTRY, neither significantly affects overall runtime. Specifically, FORCLIFT compilation typically takes around 0.5s, while GANTRY compilation takes around 2.3s.

Based on our experiments, which algorithm should be used in practice? If the formula can be handled by FORCLIFT and the domain sizes are reasonably small, FORCLIFT is likely the fastest algorithm. In other situations, GANTRY is expected to be significantly more

efficient than FASTWFOMC regardless of domain size, provided both algorithms can handle the formula.

5 Conclusion and Future Work

In this work, we have presented a scalable automated FOKC-based approach to FOMC. Our algorithm involves completing the definitions of recursive functions and subsequently translating all function definitions into C++ code. Empirical results demonstrate that GANTRY can scale to larger domain sizes than FASTWFOMC while supporting a wider range of formulas than FORCLIFT. The ability to efficiently handle large domain sizes is particularly crucial in the weighted setting, as illustrated by the ‘friends’ example discussed in Section 4, where the model captures complex social networks with probabilistic relationships. Without this scalability, the practical usefulness of these models would be limited.

Future directions for research include conducting a comprehensive experimental comparison of FOMC algorithms to better understand their comparative performance across various formulas. The capabilities of GANTRY could also be characterised theoretically, e.g. by proving completeness for specific logic fragments like C^2 . Additionally, the efficiency of FOMC algorithms can be further analysed using fine-grained complexity, which would provide more detailed insights into the computational demands of different formulas.

References

- 1 Damiano Azzolini and Fabrizio Riguzzi. Lifted inference for statistical statements in probabilistic answer set programming. *Int. J. Approx. Reason.*, 163:109040, 2023.
- 2 Paul Beame, Guy Van den Broeck, Eric Gribkoff, and Dan Suciu. Symmetric weighted first-order model counting. In *PODS*, pages 313–328. ACM, 2015.
- 3 Mark Chavira and Adnan Darwiche. On probabilistic inference by weighted model counting. *Artif. Intell.*, 172(6-7):772–799, 2008.
- 4 Adnan Darwiche. On the tractable counting of theory models and its application to truth maintenance and belief revision. *Journal of Applied Non-Classical Logics*, 11(1-2):11–34, 2001.
- 5 Paulius Dilkas and Vaishak Belle. Synthesising recursive functions for first-order model counting: Challenges, progress, and conjectures. In *KR*, pages 198–207, 2023.
- 6 Vibhav Gogate and Pedro M. Domingos. Probabilistic theorem proving. *Commun. ACM*, 59(7):107–115, 2016.
- 7 Eric Gribkoff, Dan Suciu, and Guy Van den Broeck. Lifted probabilistic inference: A guide for the database researcher. *IEEE Data Eng. Bull.*, 37(3):6–17, 2014.
- 8 Peter G. Hinman. *Fundamentals of mathematical logic*. CRC Press, 2018.
- 9 Wilfrid Hodges. *A Shorter Model Theory*. Cambridge University Press, 1997.
- 10 Manfred Jaeger and Guy Van den Broeck. Liftability of probabilistic inference: Upper and lower bounds. In *StarAI@UAI*, 2012.
- 11 Seyed Mehran Kazemi and David Poole. Knowledge compilation for lifted probabilistic inference: Compiling to a low-level language. In *KR*, pages 561–564. AAAI Press, 2016.
- 12 Kristian Kersting. Lifted probabilistic inference. In *ECAI*, volume 242 of *Frontiers in Artificial Intelligence and Applications*, pages 33–38. IOS Press, 2012.
- 13 Ondrej Kuželka. Weighted first-order model counting in the two-variable fragment with counting quantifiers. *J. Artif. Intell. Res.*, 70:1281–1307, 2021.
- 14 Sagar Malhotra and Luciano Serafini. Weighted model counting in FO2 with cardinality constraints and counting quantifiers: A closed form formula. In *AAAI*, pages 5817–5824. AAAI Press, 2022.
- 15 Nils J. Nilsson. Probabilistic logic. *Artif. Intell.*, 28(1):71–87, 1986.

- 412 16 Feng Niu, Christopher Ré, AnHai Doan, and Jude W. Shavlik. Tuffy: Scaling up statistical
413 inference in Markov logic networks using an RDBMS. *Proc. VLDB Endow.*, 4(6):373–384,
414 2011.
- 415 17 Vilém Novák, Irina Perfilieva, and Jiri Mockor. *Mathematical principles of fuzzy logic*, volume
416 517. Springer Science & Business Media, 2012.
- 417 18 Fabrizio Riguzzi, Elena Bellodi, Riccardo Zese, Giuseppe Cota, and Evelina Lamma. A
418 survey of lifted inference approaches for probabilistic logic programming under the distribution
419 semantics. *Int. J. Approx. Reason.*, 80:313–333, 2017.
- 420 19 Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach (4th Edition)*.
421 Pearson, 2020.
- 422 20 Dragan Z. Šaletić. Graded logics. *Interdisciplinary Description of Complex Systems: INDECS*,
423 22(3):276–295, 2024.
- 424 21 Parag Singla and Pedro M. Domingos. Lifted first-order belief propagation. In *AAAI*, pages
425 1094–1099. AAAI Press, 2008.
- 426 22 Martin Svatos, Peter Jung, Jan Tóth, Yuyi Wang, and Ondrej Kuželka. On discovering
427 interesting combinatorial integer sequences. In *IJCAI*, pages 3338–3346. ijcai.org, 2023.
- 428 23 Jan Tóth and Ondrej Kuželka. Lifted inference with linear order axiom. In *AAAI*, pages
429 12295–12304. AAAI Press, 2023.
- 430 24 Pietro Totis, Jesse Davis, Luc De Raedt, and Angelika Kimmig. Lifted reasoning for combina-
431 torial counting. *J. Artif. Intell. Res.*, 76:1–58, 2023.
- 432 25 Jan Tóth and Ondřej Kuželka. Complexity of weighted first-order model counting in the
433 two-variable fragment with counting quantifiers: A bound to beat, 2024. URL: <https://arxiv.org/abs/2404.12905>, arXiv:2404.12905.
- 434 26 Timothy van Bremen and Ondrej Kuželka. Approximate weighted first-order model counting:
435 Exploiting fast approximate model counters and symmetry. In *IJCAI*, pages 4252–4258.
436 ijcai.org, 2020.
- 437 27 Timothy van Bremen and Ondrej Kuželka. Faster lifting for two-variable logic using cell
438 graphs. In *UAI*, volume 161 of *Proceedings of Machine Learning Research*, pages 1393–1402.
439 AUAI Press, 2021.
- 440 28 Timothy van Bremen and Ondrej Kuželka. Lifted inference with tree axioms. *Artif. Intell.*,
441 324:103997, 2023.
- 442 29 Guy Van den Broeck. On the completeness of first-order knowledge compilation for lifted
443 probabilistic inference. In *NIPS*, pages 1386–1394, 2011.
- 444 30 Guy Van den Broeck, Arthur Choi, and Adnan Darwiche. Lifted relax, compensate and then
445 recover: From approximate to exact lifted probabilistic inference. In *UAI*, pages 131–141.
446 AUAI Press, 2012.
- 447 31 Guy Van den Broeck, Wannes Meert, and Adnan Darwiche. Skolemization for weighted
448 first-order model counting. In *KR*. AAAI Press, 2014.
- 449 32 Guy Van den Broeck, Nima Taghipour, Wannes Meert, Jesse Davis, and Luc De Raedt. Lifted
450 probabilistic inference by first-order knowledge compilation. In *IJCAI*, pages 2178–2185.
451 IJCAI/AAAI, 2011.
- 452 33 Yuanhong Wang, Juhua Pu, Yuyi Wang, and Ondrej Kuželka. Lifted algorithms for symmetric
453 weighted first-order model sampling. *Artif. Intell.*, 331:104114, 2024.
- 454