

Towards Practical First-Order Model Counting

Ananth K. Kidambi ✉

Indian Institute of Technology Bombay, Mumbai, India

Guramrit Singh ✉

Indian Institute of Technology Bombay, Mumbai, India

Paulius Dilkas ✉🏠

University of Toronto, Toronto, Canada

Vector Institute, Toronto, Canada

Kuldeep S. Meel ✉🏠

University of Toronto, Toronto, Canada

Abstract

First-order model counting (FOMC) is the problem of counting the number of models of a sentence in first-order logic. Since lifted inference techniques rely on reductions to variants of FOMC, the design of scalable methods for FOMC has attracted attention from both theoreticians and practitioners over the past decade. Recently, a new approach based on first-order knowledge compilation was proposed. This approach, called CRANE, instead of simply providing the final count, generates definitions of (possibly recursive) functions that can be evaluated with different arguments to compute the model count for any domain size. However, this approach is not fully automated, as it requires manual evaluation of the constructed functions. The primary contribution of this work is a fully automated compilation algorithm, called GANTRY, which transforms the function definitions into C++ code equipped with arbitrary-precision arithmetic. These additions allow the new FOMC algorithm to scale to domain sizes over 500,000 times larger than the current state of the art, as demonstrated through experimental results.

2012 ACM Subject Classification Theory of computation → Automated reasoning; Theory of computation → Logic and verification; Mathematics of computing → Combinatorics

Keywords and phrases first-order model counting, knowledge compilation, lifted inference

Funding Ananth K. Kidambi: TODO: funding

Guramrit Singh: TODO: funding

Paulius Dilkas: TODO: funding

Kuldeep S. Meel: TODO: funding

1 Introduction

First-order model counting (FOMC) is the task of determining the number of models for a sentence in first-order logic over a specified domain. The weighted variant, WFOMC, computes the total weight of these models, linking logical reasoning with probabilistic frameworks [32]. It builds upon earlier efforts in weighted model counting for propositional logic [3] and broader attempts to bridge logic and probability [15, 17, 20]. WFOMC is central to *lifted inference*, which enhances the efficiency of probabilistic calculations by exploiting symmetries [12]. Lifted inference continues to advance, with applications extending to constraint satisfaction problems [24] and probabilistic answer set programming [1]. Moreover, WFOMC has proven effective at reasoning over probabilistic databases [7] and probabilistic logic programs [18]. FOMC algorithms have also facilitated breakthroughs in discovering integer sequences [22] and developing recurrence relations for these sequences [5]. Recently, these algorithms have been extended to perform sampling tasks [33].

The complexity of FOMC is generally measured by *data complexity*, with a formula classified as *liftable* if it can be solved in polynomial time relative to the domain size [10].

While all formulas with up to two variables are known to be liftable [29, 31], Beame et al. [2] demonstrated that liftability does not extend to all formulas, identifying an unliftable formula with three variables. Recent work has further extended the liftable fragment with additional axioms [23, 28] and counting quantifiers [13], expanding our understanding of liftability.

FOMC algorithms are diverse, with approaches ranging from *first-order knowledge compilation* (FOKC) to local search [16], Monte Carlo sampling [6], and anytime approximation [26]. Among these, FOKC-based algorithms are particularly prominent, transforming formulas into structured representations such as circuits or graphs. Notable examples include FORCLIFT [32] and its successor CRANE [5]. Another important algorithm, FASTWFOMC [27], uses cell enumeration as its foundation.

The CRANE algorithm marked a significant step forward, expanding the range of formulas handled by FOMC algorithms. However, it had notable limitations:

1. it required manual evaluation of function definitions to compute model counts, and
2. it introduced recursive functions without proper base cases, making it more complex to use.

To address these shortcomings, we present GANTRY, a fully automated FOMC algorithm that overcomes the constraints of its predecessor. GANTRY can handle domain sizes over 500,000 times larger than previous algorithms and simplifies the user experience by automatically handling base cases and compiling function definitions into efficient C++ programs.

In Section 2, we cover some preliminaries, and in Section 3, we detail all our technical contributions. Finally, in Section 4, we present our experimental results, demonstrating GANTRY’s performance compared to other FOMC algorithms, and, in Section 5, we conclude the paper by discussing promising avenues for future work.

2 Preliminaries

In Section 2.1, we summarise the basic principles of first-order logic. Then, in Section 2.2, we formally define (W)FOMC and discuss the distinctions between three variations of first-order logic used for FOMC. Finally, in Section 2.3, we introduce the terminology used to describe the output of the original CRANE algorithm, i.e., functions and equations that define them.

We use \mathbb{N}_0 to represent the set of non-negative integers. In both algebra and logic, we write $S\sigma$ to denote the application of a *substitution* σ to an expression S , where $\sigma = [x_1 \mapsto y_1, x_2 \mapsto y_2, \dots, x_n \mapsto y_n]$ signifies the replacement of all instances of x_i with y_i for all $i = 1, \dots, n$.

2.1 First-Order Logic

In this section, we will review the basic concepts of first-order logic as they are used in FOKC algorithms. There are two key differences between the logic used by these algorithms and the logic supported as input. First, Skolemization [31] eliminates existential quantifiers by introducing additional predicates. Please note that Skolemization here differs from the standard Skolemization procedure that introduces function symbols [9]. Second, the input formula is rewritten as a conjunction of clauses, each in *prenex normal form* [8].

A *term* can be either a variable or a constant. An *atom* can be either

1. $P(t_1, \dots, t_m)$ for some predicate P and terms t_1, \dots, t_m (written as $P(\mathbf{t})$ for short) or
2. $x = y$ for some terms x and y .

The *arity* of a predicate is the number of arguments it takes, i.e., m in the case of the predicate P mentioned above. We write P/m to denote a predicate along with its arity. A *literal* can be either an atom (i.e., a *positive literal*) or its negation (i.e., a *negative literal*).

Logic	Sorts	Constants	Variables	Quantifiers	Additional atoms
FO	one or more	✓	unlimited	\forall, \exists	$x = y$
C^2	one	✗	two	$\forall, \exists, \exists^{=k}, \exists^{\leq k}, \exists^{\geq k}$	—
$UFO^2 + CC$	one	✗	two	\forall	$ P = m$

■ **Table 1** A comparison of the three logics used in FOMC based on the following aspects:

1. the number of sorts,
2. support for constants,
3. the maximum number of variables,
4. supported quantifiers, and
5. supported atoms in addition to those of the form $P(\mathbf{t})$ for a predicate P/n and an n -tuple of terms \mathbf{t} .

Here:

1. k and m are non-negative integers, with the latter depending on the domain size,
2. P represents a predicate, and
3. x and y are terms.

91 An atom is *ground* if it contains no variables, i.e., only constants. A *clause* is of the form
 92 $\forall x_1 \in \Delta_1. \forall x_2 \in \Delta_2 \dots \forall x_n \in \Delta_n. \phi(x_1, x_2, \dots, x_n)$, where ϕ is a disjunction of literals that
 93 only contain variables x_1, \dots, x_n (and any constants). We say that a clause is a (*positive*)
 94 *unit clause* if

- 95 1. there is only one literal with a predicate, and
- 96 2. it is a positive literal.

97 Finally, a *formula* is a conjunction of clauses. Throughout the paper, we will use set-theoretic
 98 notation, interpreting a formula as a set of clauses and a clause as a set of literals.

99 2.2 FOMC Algorithms and Their Logics

100 In Table 1, we outline the differences among three first-order logics commonly used in FOMC:

- 101 1. FO is the input format for FORCLIFT* and its extensions CRANE† and GANTRY;
- 102 2. C^2 is often used in the literature on FASTWFOMC and related methods [13, 14];
- 103 3. $UFO^2 + CC$ is the input format supported by the most recent implementation of FAST-
 104 WFOMC‡.

105 The notation we use to refer to each logic is standard in the case of C^2 and $UFO^2 + CC$ [25] and
 106 redefined to be more specific in the case of FO. All three logics are function-free, and domains
 107 are always assumed to be finite. As usual, we presuppose the *unique name assumption*, which
 108 states that two constants are equal if and only if they are the same constant [19].

109 In FO, each term is assigned to a *sort*, and each predicate P/n is assigned to a sequence
 110 of n sorts. Each sort has its corresponding domain. These assignments to sorts are typically
 111 left implicit and can be reconstructed from the quantifiers. For example, $\forall x, y \in \Delta. P(x, y)$
 112 implies that variables x and y have the same sort. On the other hand, $\forall x \in \Delta. \forall y \in \Gamma. P(x, y)$
 113 implies that x and y have different sorts, and it would be improper to write, for example,

* <https://github.com/UCLA-StarAI/Forclift>

† <https://doi.org/10.5281/zenodo.8004077>

‡ <https://github.com/jan-toth/FastWFOMC.jl>

114 $\forall x \in \Delta. \forall y \in \Gamma. P(x, y) \vee x = y$. FO is also the only logic to support constants, formulas
 115 with more than two variables, and the equality predicate. While we do not explicitly refer to
 116 sorts in subsequent sections of this paper, the many-sorted nature of FO is paramount to the
 117 algorithms presented therein.

118 ► **Remark 1.** In the case of FORCLIFT and its extensions, support for a formula as valid input
 119 does not imply that the algorithm can compile the formula into a circuit or graph suitable
 120 for lifted model counting. However, it is known that FORCLIFT compilation is guaranteed to
 121 succeed on any FO formula without constants and with at most two variables [29, 31].

122 Compared to FO, C^2 and $UFO^2 + CC$ lack support for

- 123 1. constants,
- 124 2. the equality predicate,
- 125 3. multiple domains, and
- 126 4. formulas with more than two variables.

127 The advantage that C^2 brings over FO is the inclusion of *counting quantifiers*. That is,
 128 alongside \forall and \exists , C^2 supports $\exists^{=k}$, $\exists^{\leq k}$, and $\exists^{\geq k}$ for any positive integer k . For example,
 129 $\exists^{=1}x. \phi(x)$ means that there exists *exactly one* x such that $\phi(x)$, and $\exists^{\leq 2}x. \phi(x)$ means
 130 that there exist *at most two* such x . $UFO^2 + CC$, on the other hand, does not support any
 131 existential quantifiers but instead incorporates (*equality*) *cardinality constraints*. For example,
 132 $|P| = 3$ constrains all models to have *precisely three positive literals with the predicate* P .

133 ► **Definition 2 (Model).** Let ϕ be a formula in FO. For each predicate P/n in ϕ , let $(\Delta_i^P)_{i=1}^n$
 134 be a list of the corresponding domains. Let σ be a map from the domains of ϕ to their
 135 interpretations as sets, satisfying the following conditions:

- 136 1. the sets are pairwise disjoint, and
- 137 2. the constants in ϕ are included in the corresponding domains.

138 A structure of ϕ is a set M of ground literals defined by adding to M either $P(\mathbf{t})$ or $\neg P(\mathbf{t})$
 139 for every predicate P/n in ϕ and n -tuple $\mathbf{t} \in \prod_{i=1}^n \sigma(\Delta_i^P)$. A structure is a model if it
 140 satisfies ϕ .

141 ► **Remark 3.** The distinctness of domains is important in two ways. First, in terms of
 142 expressiveness, a clause such as $\forall x \in \Delta. P(x, x)$ is valid if predicate P is defined over two
 143 copies of the same domain and invalid otherwise. Second, having more distinct domains
 144 makes the problem more decomposable for the FOKC algorithm. With distinct domains, the
 145 algorithm can make assumptions or deductions about, e.g., the first domain of predicate P
 146 without worrying how (or if) they apply to the second domain.

147 While this work focuses on FOMC, we still define the weighted variant of the problem as
 148 Skolemization relies on weights even for unweighted FOMC.

149 ► **Definition 4 (WFOMC instance).** A WFOMC instance *comprises*:

- 150 1. a formula ϕ in FO,
 - 151 2. two (rational) weights $w^+(P)$ and $w^-(P)$ assigned to each predicate P in ϕ , and
 - 152 3. σ as described in Definition 2.
- 153 Unless specified otherwise, we assume all weights to be equal to 1.

154 ► **Definition 5 (WFOMC [32]).** Given a WFOMC instance (ϕ, w^+, w^-, σ) as in Definition 4,
 155 the (symmetric) weighted first-order model count (WFOMC) of ϕ is

$$156 \sum_{M \models \phi} \prod_{P(\mathbf{t}) \in M} w^+(P) \prod_{\neg P(\mathbf{t}) \in M} w^-(P), \quad (1)$$

157 where the sum is over all models of ϕ .

► **Example 6** (Counting functions). To define predicate P as a function from a domain Δ to itself, in \mathcal{C}^2 one would write $\forall x \in \Delta. \exists^{=1} y \in \Delta. P(x, y)$. In $\text{UFO}^2 + \text{CC}$, the same could be written as

$$(\forall x, y \in \Delta. S(x) \vee \neg P(x, y)) \wedge (|P| = |\Delta|), \quad (2)$$

where $w^-(S) = -1$. Although Formula (2) has more models compared to its counterpart in \mathcal{C}^2 , the negative weight $w^-(S) = -1$ makes some of the terms in Equation (1) cancel out.

Equivalently, in FO we would write

$$(\forall x \in \Gamma. \exists y \in \Delta. P(x, y)) \wedge (\forall x \in \Gamma. \forall y, z \in \Delta. P(x, y) \wedge P(x, z) \Rightarrow y = z). \quad (3)$$

The first clause asserts that each x must have at least one corresponding y , while the second statement adds the condition that if x is mapped to both y and z , then y must equal z . It is important to note that Formula (3) is written with two domains instead of just one. However, we can still determine the correct number of functions by assuming that the sizes of Γ and Δ are equal. This formulation, as observed by Dilkas and Belle [5], can prove beneficial in enabling FOKC algorithms to find efficient solutions.

2.3 Algebra

We write **expr** to represent an arbitrary algebraic expression. It is important to note that some terms have different meanings in algebra and logic. In algebra, a *constant* refers to a non-negative integer. Likewise, a *variable* can either be a parameter of a function or a variable introduced through summation, such as i in the expression $\sum_{i=1}^n \text{expr}$. A (function) *signature* is $f(x_1, \dots, x_n)$ (or $f(\mathbf{x})$ for short), where f represents an n -ary function, and each x_i represents a variable. An *equation* is $f(\mathbf{x}) = \text{expr}$, with $f(\mathbf{x})$ representing a signature.

► **Definition 7** (Base case). Let $f(\mathbf{x})$ be a function call where each x_i is either a constant or a variable (note that signatures are included in this definition). Then function call $f(\mathbf{y})$ is considered a base case of $f(\mathbf{x})$ if $f(\mathbf{y}) = f(\mathbf{x})\sigma$, where σ is a substitution that replaces one or more x_i with a constant.

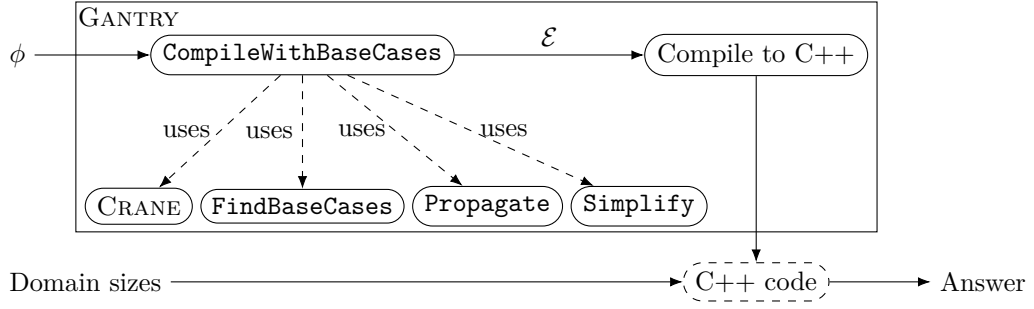
3 Technical Contributions

Figure 1 provides an overview of GANTRY’s workflow. Section 3.1 describes the main algorithm for completing the definitions of recursive functions with a sufficient set of base cases. Sections 3.2 and 3.3 describe subsidiary algorithms for constructing a set of base cases and their corresponding logical formulas. Section 3.4 explains the post-processing techniques for ensuring accurate model counting. Additionally, Section 3.5 explains the process of compiling function definitions into C++ code, greatly expanding upon the range of formulas that could previously be handled by similar approaches [11].

3.1 Completing the Definitions of Functions

Before describing the main contribution of this work, let us review the essential aspects of FOKC as realised by CRANE. The input formula is compiled into:

1. set \mathcal{E} of equations,
2. map \mathcal{F} from function names to formulas, and
3. map \mathcal{D} from function names and argument indices to domains.



■ **Figure 1** The outline of using GANTRY to compute the model count of a formula ϕ . First, the formula is compiled into a set of equations, which are then used to create a C++ program. This program can be executed with different command line arguments to calculate the model count of ϕ for different domain sizes. To accomplish this, the `CompileWithBaseCases` function employs several components:

1. the FOKC algorithm of `CRANE`,
2. a procedure called `FindBaseCases`, which identifies a sufficient set of base cases,
3. a procedure called `Propagate`, which constructs a formula corresponding to a given base case, and
4. algebraic simplification techniques (denoted as `Simplify`).

■ **Algorithm 1** `CompileWithBaseCases(ϕ)`

Input: formula ϕ
Output: set \mathcal{E} of equations

```

1  $(\mathcal{E}, \mathcal{F}, \mathcal{D}) \leftarrow \text{CRANE}(\phi);$ 
2  $\mathcal{E} \leftarrow \text{Simplify}(\mathcal{E});$ 
3 foreach base case  $f(\mathbf{x}) \in \text{FindBaseCases}(\mathcal{E})$  do
4    $\psi \leftarrow \mathcal{F}(f);$ 
5   foreach index  $i$  such that  $x_i \in \mathbb{N}_0$  do  $\psi \leftarrow \text{Propagate}(\psi, \mathcal{D}(f, i), x_i);$ 
6    $\mathcal{E} \leftarrow \mathcal{E} \cup \text{CompileWithBaseCases}(\psi);$ 

```

\mathcal{E} can contain any number of functions, one of which (denoted by f) represents the solution to the FOMC problem. To compute the FOMC for particular domain sizes, f must be evaluated with those domain sizes as arguments. \mathcal{D} records this correspondence between function arguments and domains.

Algorithm 1 presents our overall approach for compiling a formula into equations that include the necessary base cases. To begin, we use the FOKC algorithm of the original `CRANE` to compile the formula into the three components: \mathcal{E} , \mathcal{F} , and \mathcal{D} . After some algebraic simplification, \mathcal{E} is passed to the `FindBaseCases` procedure (see Section 3.2). For each base case $f(\mathbf{x})$, we retrieve the logical formula $\mathcal{F}(f)$ associated with the function name f and simplify it using the `Propagate` procedure (explained in detail in Section 3.3). We do this by iterating over all indices of \mathbf{x} , where x_i is a constant, and using `Propagate` to simplify ψ by assuming that domain $\mathcal{D}(f, i)$ has size x_i . Finally, on line 6, `CompileWithBaseCases` recurses on these simplified formulas and adds the resulting base case equations to \mathcal{E} . Example 9 below provides more detail.

► **Remark 8.** Although `CompileWithBaseCases` starts with a call to `CRANE`, the proposed algorithm is not just a post-processing step for FOKC because Algorithm 1 is recursive and

213 can issue more calls to CRANE on various derived formulas.

214 ► **Example 9** (Counting bijections). Consider the following formula (previously examined by
215 Dilkas and Belle [5]) that defines predicate P as a bijection between two sets Γ and Δ :

$$\begin{aligned} & (\forall x \in \Gamma. \exists y \in \Delta. P(x, y)) \wedge \\ & (\forall y \in \Delta. \exists x \in \Gamma. P(x, y)) \wedge \\ 216 & (\forall x \in \Gamma. \forall y, z \in \Delta. P(x, y) \wedge P(x, z) \Rightarrow y = z) \wedge \\ & (\forall x, z \in \Gamma. \forall y \in \Delta. P(x, y) \wedge P(z, y) \Rightarrow x = z). \end{aligned}$$

217 We specifically examine the first solution returned by GANTRY for this formula.

218 After line 2, we have

$$\begin{aligned} 219 \quad \mathcal{E} &= \left\{ \begin{array}{l} f(m, n) = \sum_{l=0}^n \binom{n}{l} (-1)^{n-l} g(l, m), \\ g(l, m) = g(l-1, m) + mg(l-1, m-1) \end{array} \right\}; \\ 220 \quad \mathcal{D} &= \{ (f, 1) \mapsto \Gamma, (f, 2) \mapsto \Delta, (g, 1) \mapsto \Delta^\top, (g, 2) \mapsto \Gamma \}, \end{aligned}$$

221 where Δ^\top is a new domain. (We omit the definition of \mathcal{F} as the formulas can get a bit
222 verbose.) Then **FindBaseCases** identifies two base cases: $g(0, m)$ and $g(l, 0)$. In both cases,
223 **CompileWithBaseCases** recurses on the formula $\mathcal{F}(g)$ simplified by assuming that one of the
224 domains is empty. In the first case, we recurse on the formula $\forall x \in \Gamma. S(x) \vee \neg S(x)$, where
225 S is a predicate introduced by Skolemization with weights $w^+(S) = 1$ and $w^-(S) = -1$.
226 Hence, we obtain the base case $g(0, m) = 0^m$. In the case of $g(l, 0)$, **Propagate**($\psi, \Gamma, 0$)
227 returns an empty formula, resulting in $g(l, 0) = 1$.

228 It is worth noting that these base cases overlap when $l = m = 0$ but remain consistent
229 since $0^0 = 1$. Generally, let ϕ be a formula with two domains Γ and Δ , and let $n, m \in \mathbb{N}_0$.
230 Then the FOMC of **Propagate**(ϕ, Δ, n) assuming $|\Gamma| = m$ is the same as the FOMC of
231 **Propagate**(ϕ, Γ, m) assuming $|\Delta| = n$.

232 Finally, the main responsibility of the **Simplify** procedure is to handle the algebraic
233 pattern $\sum_{m=0}^n [a \leq m \leq b] f(m)$. Here:

- 234 1. n is a variable,
- 235 2. $a, b \in \mathbb{N}_0$ are constants,
- 236 3. f is an expression that may depend on m , and
- 237 4. $[a \leq m \leq b] = \begin{cases} 1 & \text{if } a \leq m \leq b \\ 0 & \text{otherwise} \end{cases}$.

238 **Simplify** transforms this pattern into $f(a) + f(a+1) + \dots + f(\min\{n, b\})$. For instance,
239 in the case of Example 9, **Simplify** transforms $g(l, m) = \sum_{k=0}^m [0 \leq k \leq 1] \binom{m}{k} g(l-1, m-k)$
240 into $g(l, m) = g(l-1, m) + mg(l-1, m-1)$.

241 3.2 Identifying a Sufficient Set of Base Cases

242 Algorithm 2 summarises the implementation of **FindBaseCases**. **FindBaseCases** considers
243 two types of arguments when a function f calls itself recursively:

- 244 1. constants and
- 245 2. arguments of the form $x_i - c_i$, where c_i is a constant and x_i is the i -th argument of the
246 signature of f .

247 When the argument is a constant c_i , a base case with c_i is added. In the second case, a base
248 case is added for each constant from 0 up to (but not including) c_i .

Algorithm 2 FindBaseCases(\mathcal{E})

Input: set \mathcal{E} of equations
Output: set \mathcal{B} of base cases
1 $\mathcal{B} \leftarrow \emptyset$;
2 **foreach** function call $f(\mathbf{y})$ on the right-hand side of an equation in \mathcal{E} **do**
3 $\mathbf{x} \leftarrow$ the parameters of f in its definition;
4 **foreach** $y_i \in \mathbf{y}$ **do**
5 **if** $y_i \in \mathbb{N}_0$ **then** $\mathcal{B} \leftarrow \mathcal{B} \cup \{f(\mathbf{x})[x_i \mapsto y_i]\}$;
6 **else if** $y_i = x_i - c_i$ for some $c_i \in \mathbb{N}_0$ **then**
7 **for** $j \leftarrow 0$ **to** $c_i - 1$ **do** $\mathcal{B} \leftarrow \mathcal{B} \cup \{f(\mathbf{x})[x_i \mapsto j]\}$;

249 **► Example 10.** Consider the recursive function g from Example 9. FindBaseCases iterates
250 over two function calls: $g(l-1, m)$ and $g(l-1, m-1)$. The former produces the base case
251 $g(0, m)$, while the latter produces both $g(0, m)$ and $g(l, 0)$.

252 It can be shown that the base cases identified by FindBaseCases are sufficient for the
253 algorithm to terminate.⁴

254 **► Theorem 11 (Termination).** Let \mathcal{E} represent the equations returned by CompileWithBaseCases.
255 Let f be an n -ary function in \mathcal{E} and $\mathbf{x} \in \mathbb{N}_0^n$. Then the evaluation of $f(\mathbf{x})$ terminates.

256 We prove Theorem 11 using double induction. First, we apply induction to the number
257 of functions in \mathcal{E} . Then, we use induction on the arity of the ‘last’ function in \mathcal{E} according to
258 some topological ordering. For the detailed proof, please refer to the technical appendix.

259 3.3 Propagating Domain Size Assumptions

260 Algorithm 3, called Propagate, modifies the formula ϕ based on the assumption that $|\Delta| = n$.
261 When $n = 0$, some clauses become vacuously satisfied and can be removed. When $n > 0$,
262 partial grounding is performed by replacing all variables quantified over Δ with constants.
263 (None of the formulas examined in this work had $n > 1$.) Algorithm 3 handles these two
264 cases separately. For a literal or a clause C , the set of corresponding domains is denoted as
265 Doms(C).

266 In the case of $n = 0$, there are three types of clauses to consider:

- 267 1. those that do not mention Δ ,
- 268 2. those in which every literal contains variables quantified over Δ , and
- 269 3. those that have some literals with variables quantified over Δ and some without.

270 Clauses of Type 1 are transferred to the new formula ϕ' without any changes. For clauses of
271 Type 2, C' is empty, so these clauses are filtered out. As for clauses of Type 3, a new kind of
272 smoothing is performed, which will be explained in Section 3.4.

273 In the case of $n > 0$, n new constants are introduced. Let C be an arbitrary clause in ϕ ,
274 and let $m \in \mathbb{N}_0$ be the number of variables in C quantified over Δ . If $m = 0$, C is added
275 directly to ϕ' . Otherwise, a clause is added to ϕ' for every possible combination of replacing
276 the m variables in C with the n new constants.

⁴ Note that characterising the fine-grained complexity of the solutions found by GANTRY or other FOMC algorithms is an emerging area of research. These questions have been partially addressed in previous work [5, 25] and are orthogonal to the goals of this section.

Algorithm 3 $\text{Propagate}(\phi, \Delta, n)$

Input: formula ϕ , domain Δ , $n \in \mathbb{N}_0$
Output: formula ϕ'

```

1  $\phi' \leftarrow \emptyset$ ;
2 if  $n = 0$  then
3   foreach  $\text{clause } C \in \phi$  do
4     if  $\Delta \notin \text{Doms}(C)$  then  $\phi' \leftarrow \phi' \cup \{C\}$ ;
5     else
6        $C' \leftarrow \{l \in C \mid \Delta \notin \text{Doms}(l)\}$ ;
7       if  $C' \neq \emptyset$  then
8          $l \leftarrow \text{an arbitrary literal in } C'$ ;
9          $\phi' \leftarrow \phi' \cup \{C' \cup \{\neg l\}\}$ ;
10 else
11    $D \leftarrow \text{a set of } n \text{ new constants in } \Delta$ ;
12   foreach  $\text{clause } C \in \phi$  do
13      $(x_i)_{i=1}^m \leftarrow \text{the variables in } C \text{ with domain } \Delta$ ;
14     if  $m = 0$  then  $\phi' \leftarrow \phi' \cup \{C\}$ ;
15     else  $\phi' \leftarrow \phi' \cup \{C[x_1 \mapsto c_1, \dots, x_m \mapsto c_m] \mid (c_i)_{i=1}^m \in D^m\}$ ;

```

277 **► Example 12.** Let $C \equiv \forall x \in \Gamma. \forall y, z \in \Delta. \neg P(x, y) \vee \neg P(x, z) \vee y = z$. Then $\text{Doms}(C) =$
 278 $\text{Doms}(\neg P(x, y)) = \text{Doms}(\neg P(x, z)) = \{\Gamma, \Delta\}$, and $\text{Doms}(y = z) = \{\Delta\}$. A call to
 279 $\text{Propagate}(\{C\}, \Delta, 3)$ would result in the following formula with nine clauses:

$$\begin{aligned}
 & (\forall x \in \Gamma. \neg P(x, c_1) \vee \neg P(x, c_1) \vee c_1 = c_1) \wedge \\
 & (\forall x \in \Gamma. \neg P(x, c_1) \vee \neg P(x, c_2) \vee c_1 = c_2) \wedge \\
 & \quad \vdots \\
 & (\forall x \in \Gamma. \neg P(x, c_3) \vee \neg P(x, c_3) \vee c_3 = c_3).
 \end{aligned}$$

284 Here, c_1 , c_2 , and c_3 are the new constants.

285 3.4 Smoothing the Base Cases

286 *Smoothing* modifies a circuit to reintroduce eliminated atoms, ensuring the correct model
 287 count [4, 32]. In this section, we describe a similar process performed on lines 7–9 of
 288 Algorithm 3. Line 7 checks if smoothing is necessary, and lines 8 and 9 execute it. If the
 289 condition on line 7 is not satisfied, the clause is not smoothed but omitted.

290 Suppose Propagate is called with arguments $(\phi, \Delta, 0)$, i.e., we are simplifying the formula
 291 ϕ by assuming that the domain Δ is empty. Informally, if there is a predicate P in ϕ unrelated
 292 to Δ , smoothing preserves all occurrences of P even if all clauses with P become vacuously
 293 satisfied.

294 **► Example 13.** Let ϕ be

$$(\forall x \in \Delta. \forall y, z \in \Gamma. Q(x) \vee P(y, z)) \wedge \quad (4)$$

$$(\forall y, z \in \Gamma'. P(y, z)), \quad (5)$$

where $\Gamma' \subseteq \Gamma$ is a domain introduced by a compilation rule. It should be noted that P , as a relation, is a subset of $\Gamma \times \Gamma$.

Now, let us reason manually about the model count of ϕ when $\Delta = \emptyset$. Predicate Q can only take one value, $Q = \emptyset$. The value of P is fixed over $\Gamma' \times \Gamma'$ by Clause (5), but it can vary freely over $(\Gamma \times \Gamma) \setminus (\Gamma' \times \Gamma')$ since Clause (4) is vacuously satisfied by all structures. Therefore, the correct FOMC should be $2^{|\Gamma|^2 - |\Gamma'|^2}$. However, without line 9, **Propagate** would simplify ϕ to $\forall y, z \in \Gamma'. P(y, z)$. In this case, P is a subset of $\Gamma' \times \Gamma'$. This simplified formula has only one model: $\{P(y, z) \mid y, z \in \Gamma'\}$. By including line 9, **Propagate** transforms ϕ to

$$(\forall y, z \in \Gamma. P(y, z) \vee \neg P(y, z)) \wedge (\forall y, z \in \Gamma'. P(y, z)),$$

which retains the correct model count.

It is worth mentioning that the choice of l on line 8 of Algorithm 3 is inconsequential because any choice achieves the same goal: constructing a tautological clause that retains the literals in C' .

3.5 Generating C++ Code

In this section, we will describe the final step of GANTRY as outlined in Figure 1. This step involves translating the set of equations \mathcal{E} into C++ code. The resulting C++ program can then be compiled and executed with different command-line arguments to compute the model count of the formula for various domain sizes.

Each equation in \mathcal{E} is compiled into a C++ function, along with a separate cache for memoisation. Let us consider an arbitrary equation $e = (f(\mathbf{x}) = \text{expr}) \in \mathcal{E}$, and let $\mathbf{c} \in \mathbb{N}_0^n$ represent the arguments of the corresponding C++ function. The implementation of e consists of three parts. First, we check if \mathbf{c} is already present in the cache of e . If it is, we simply return the cached value. Second, for each base case $f(\mathbf{y})$ of $f(\mathbf{x})$ (as defined in Definition 7), we check if \mathbf{c} *matches* \mathbf{y} , i.e., $c_i = y_i$ whenever $y_i \in \mathbb{N}_0$. If this condition is satisfied, \mathbf{c} is redirected to the C++ function that corresponds to the definition of the base case $f(\mathbf{y})$. Finally, if none of the above cases apply, we evaluate \mathbf{c} based on the expression expr , store the result in the cache, and return it.

4 Experimental Evaluation

Our empirical evaluation sought to compare the runtime performance of GANTRY with the current state of the art, namely FASTWFOMC and FORCLIFT. It is worth remarking that FORCLIFT does not support arbitrary precision, and returns error for cases that requires arbitrary precision reasoning. Our experiments involve two versions of GANTRY: GANTRY-GREEDY and GANTRY-BFS. Like its predecessor, GANTRY has two modes for applying compilation rules to formulas: one that uses a greedy search algorithm similar to FORCLIFT and another that combines greedy and breadth-first search.

The experiments were conducted using an Intel Skylake 2.4 GHz CPU with 188 GiB of memory and CentOS 7. C++ programs were compiled using the Intel C++ Compiler 2020u4. FASTWFOMC ran on Julia 1.10.4, while the other algorithms were executed on the Java Virtual Machine 1.8.0_201.

4.1 Benchmarks

We compare these algorithms using three benchmarks from previous studies. The first benchmark is the function-counting problem from Example 6, previously examined by Dilkas




Figure 2 The runtime of the algorithms as a function of the domain size. Note that both axes are on a logarithmic scale.

and Belle [5]. The second benchmark is a variant of the well-known ‘Friends and Smokers’ Markov logic network [21, 30]. In C^2 , FO, and $UFO^2 + CC$, this problem can be formulated as

$$(\forall x, y \in \Delta. S(x) \wedge F(x, y) \Rightarrow S(y)) \wedge (\forall x \in \Delta. S(x) \Rightarrow C(x))$$

or, equivalently, in conjunctive normal form as

$$(\forall x, y \in \Delta. S(y) \vee \neg S(x) \vee \neg F(x, y)) \wedge (\forall x \in \Delta. C(x) \vee \neg S(x)).$$

Finally, we include the bijection-counting problem previously utilised by Dilkas and Belle [5].

Its formulation in FO is described in Example 9. The equivalent formula in C^2 is

$$(\forall x \in \Delta. \exists^1 y \in \Delta. P(x, y)) \wedge (\forall y \in \Delta. \exists^1 x \in \Delta. P(x, y)).$$

Similarly, in $UFO^2 + CC$ the same formula can be written as

$$(\forall x, y \in \Delta. R(x) \vee \neg P(x, y)) \wedge (\forall x, y \in \Delta. S(x) \vee \neg P(y, x)) \wedge (|P| = |\Delta|),$$

where $w^-(R) = w^-(S) = -1$.

The three benchmark families cover a wide range of possibilities. The ‘friends’ benchmark stands out as it uses multiple predicates and can be expressed in FO using just two variables without cardinality constraints or counting quantifiers. The ‘functions’ benchmark, on the other hand, can still be handled by all the algorithms, but it requires cardinality constraints, counting quantifiers, or more than two variables. Lastly, the ‘bijections’ benchmark is an example of a formula that FASTWFOMC can handle but FORCLIFT cannot.

For evaluation purposes, we ran each algorithm on each benchmark using domains of sizes $2^1, 2^2, 2^3$, and so on, until an algorithm failed to handle a domain size due to timeout, out of memory error, or out of precision errors. While we separately measured compilation and inference time, we primarily focus on total runtime, dominated by the latter.

4.2 Results

Figure 2 presents a summary of the experimental results. Only FASTWFOMC and GANTRY-BFS could handle the bijection-counting problem. For this benchmark, the largest domain sizes these algorithms could accommodate were 64 and 4096, respectively. On the other two benchmarks, FORCLIFT had the lowest runtime. However, due to its finite precision, it only scaled up to domain sizes of 16 and 128 for ‘friends’ and ‘functions’, respectively. FASTWFOMC outperformed FORCLIFT in the case of ‘friends’, but not ‘functions’, as it could handle domains of size 1024 and 64, respectively. Furthermore, both GANTRY-BFS and GANTRY-GREEDY performed similarly on both benchmarks. Similarly to the ‘bijections’ benchmark, GANTRY significantly outperformed the other two algorithms, scaling up to domains of size 8192 and 67,108,864, respectively.

Another aspect of the experimental results that deserves separate discussion is compilation. Both Julia and Scala use just-in-time (JIT) compilation, which means that FASTWFOMC and FORCLIFT take longer to run on the smallest domain size, where most JIT compilation occurs. In the case of GANTRY, it is only run once per benchmark, so the JIT compilation time is included in its overall runtime across all domain sizes. Additionally, while FORCLIFT’s compilation is generally faster than that of GANTRY, neither significantly affects overall runtime. Specifically, FORCLIFT compilation typically takes around 0.5s, while GANTRY compilation takes around 2.3s.

Based on our experiments, which algorithm should be used in practice? If the formula can be handled by FORCLIFT and the domain sizes are reasonably small, FORCLIFT is likely the fastest algorithm. In other situations, GANTRY is expected to be significantly more efficient than FASTWFOMC regardless of domain size, provided both algorithms can handle the formula.

5 Conclusion and Future Work

In this work, we have presented a scalable automated FOKC-based approach to FOMC. Our algorithm involves completing the definitions of recursive functions and subsequently translating all function definitions into C++ code. Empirical results demonstrate that GANTRY can scale to larger domain sizes than FASTWFOMC while supporting a wider range of formulas than FORCLIFT. The ability to efficiently handle large domain sizes is particularly crucial in the weighted setting, as illustrated by the ‘friends’ example discussed in Section 4, where the model captures complex social networks with probabilistic relationships. Without this scalability, the practical usefulness of these models would be limited.

Future directions for research include conducting a comprehensive experimental comparison of FOMC algorithms to better understand their comparative performance across various formulas. The capabilities of GANTRY could also be characterised theoretically, e.g. by proving completeness for specific logic fragments like C^2 . Additionally, the efficiency of FOMC algorithms can be further analysed using fine-grained complexity, which would provide more detailed insights into the computational demands of different formulas.

References

- 1 Damiano Azzolini and Fabrizio Riguzzi. Lifted inference for statistical statements in probabilistic answer set programming. *Int. J. Approx. Reason.*, 163:109040, 2023.
- 2 Paul Beame, Guy Van den Broeck, Eric Gribkoff, and Dan Suciu. Symmetric weighted first-order model counting. In *PODS*, pages 313–328. ACM, 2015.

- 404 **3** Mark Chavira and Adnan Darwiche. On probabilistic inference by weighted model counting.
405 *Artif. Intell.*, 172(6-7):772–799, 2008.
- 406 **4** Adnan Darwiche. On the tractable counting of theory models and its application to truth
407 maintenance and belief revision. *Journal of Applied Non-Classical Logics*, 11(1-2):11–34, 2001.
- 408 **5** Paulius Dilkas and Vaishak Belle. Synthesising recursive functions for first-order model
409 counting: Challenges, progress, and conjectures. In *KR*, pages 198–207, 2023.
- 410 **6** Vibhav Gogate and Pedro M. Domingos. Probabilistic theorem proving. *Commun. ACM*,
411 59(7):107–115, 2016.
- 412 **7** Eric Gribkoff, Dan Suciu, and Guy Van den Broeck. Lifted probabilistic inference: A guide
413 for the database researcher. *IEEE Data Eng. Bull.*, 37(3):6–17, 2014.
- 414 **8** Peter G. Hinman. *Fundamentals of mathematical logic*. CRC Press, 2018.
- 415 **9** Wilfrid Hodges. *A Shorter Model Theory*. Cambridge University Press, 1997.
- 416 **10** Manfred Jaeger and Guy Van den Broeck. Liftability of probabilistic inference: Upper and
417 lower bounds. In *StarAI@UAI*, 2012.
- 418 **11** Seyed Mehran Kazemi and David Poole. Knowledge compilation for lifted probabilistic
419 inference: Compiling to a low-level language. In *KR*, pages 561–564. AAAI Press, 2016.
- 420 **12** Kristian Kersting. Lifted probabilistic inference. In *ECAI*, volume 242 of *Frontiers in Artificial*
421 *Intelligence and Applications*, pages 33–38. IOS Press, 2012.
- 422 **13** Ondrej Kuželka. Weighted first-order model counting in the two-variable fragment with
423 counting quantifiers. *J. Artif. Intell. Res.*, 70:1281–1307, 2021.
- 424 **14** Sagar Malhotra and Luciano Serafini. Weighted model counting in FO2 with cardinality
425 constraints and counting quantifiers: A closed form formula. In *AAAI*, pages 5817–5824. AAAI
426 Press, 2022.
- 427 **15** Nils J. Nilsson. Probabilistic logic. *Artif. Intell.*, 28(1):71–87, 1986.
- 428 **16** Feng Niu, Christopher Ré, AnHai Doan, and Jude W. Shavlik. Tuffy: Scaling up statistical
429 inference in Markov logic networks using an RDBMS. *Proc. VLDB Endow.*, 4(6):373–384,
430 2011.
- 431 **17** Vilém Novák, Irina Perfilieva, and Jiri Mockor. *Mathematical principles of fuzzy logic*, volume
432 517. Springer Science & Business Media, 2012.
- 433 **18** Fabrizio Riguzzi, Elena Bellodi, Riccardo Zese, Giuseppe Cota, and Evelina Lamma. A
434 survey of lifted inference approaches for probabilistic logic programming under the distribution
435 semantics. *Int. J. Approx. Reason.*, 80:313–333, 2017.
- 436 **19** Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach (4th Edition)*.
437 Pearson, 2020.
- 438 **20** Dragan Z. Šaletić. Graded logics. *Interdisciplinary Description of Complex Systems: INDECS*,
439 22(3):276–295, 2024.
- 440 **21** Parag Singla and Pedro M. Domingos. Lifted first-order belief propagation. In *AAAI*, pages
441 1094–1099. AAAI Press, 2008.
- 442 **22** Martin Svatos, Peter Jung, Jan Tóth, Yuyi Wang, and Ondrej Kuželka. On discovering
443 interesting combinatorial integer sequences. In *IJCAI*, pages 3338–3346. ijcai.org, 2023.
- 444 **23** Jan Tóth and Ondrej Kuželka. Lifted inference with linear order axiom. In *AAAI*, pages
445 12295–12304. AAAI Press, 2023.
- 446 **24** Pietro Totis, Jesse Davis, Luc De Raedt, and Angelika Kimmig. Lifted reasoning for combina-
447 torial counting. *J. Artif. Intell. Res.*, 76:1–58, 2023.
- 448 **25** Jan Tóth and Ondřej Kuželka. Complexity of weighted first-order model counting in the
449 two-variable fragment with counting quantifiers: A bound to beat, 2024. URL: <https://arxiv.org/abs/2404.12905>, arXiv:2404.12905.
- 450 **26** Timothy van Bremen and Ondrej Kuželka. Approximate weighted first-order model counting:
451 Exploiting fast approximate model counters and symmetry. In *IJCAI*, pages 4252–4258.
452 ijcai.org, 2020.
- 453

- 454 **27** Timothy van Bremen and Ondrej Kuželka. Faster lifting for two-variable logic using cell
455 graphs. In *UAI*, volume 161 of *Proceedings of Machine Learning Research*, pages 1393–1402.
456 AUAI Press, 2021.
- 457 **28** Timothy van Bremen and Ondrej Kuželka. Lifted inference with tree axioms. *Artif. Intell.*,
458 324:103997, 2023.
- 459 **29** Guy Van den Broeck. On the completeness of first-order knowledge compilation for lifted
460 probabilistic inference. In *NIPS*, pages 1386–1394, 2011.
- 461 **30** Guy Van den Broeck, Arthur Choi, and Adnan Darwiche. Lifted relax, compensate and then
462 recover: From approximate to exact lifted probabilistic inference. In *UAI*, pages 131–141.
463 AUAI Press, 2012.
- 464 **31** Guy Van den Broeck, Wannes Meert, and Adnan Darwiche. Skolemization for weighted
465 first-order model counting. In *KR*. AAAI Press, 2014.
- 466 **32** Guy Van den Broeck, Nima Taghipour, Wannes Meert, Jesse Davis, and Luc De Raedt. Lifted
467 probabilistic inference by first-order knowledge compilation. In *IJCAI*, pages 2178–2185.
468 IJCAI/AAAI, 2011.
- 469 **33** Yuanhong Wang, Juhua Pu, Yuyi Wang, and Ondrej Kuželka. Lifted algorithms for symmetric
470 weighted first-order model sampling. *Artif. Intell.*, 331:104114, 2024.